

CSC 369

Operating Systems

Bogdan Simion

bogdan@cs.toronto.edu



University of Toronto, Department of Computer Science



Why are you here?

Why did the department decide that an OS course is one of only 2 required 3rd year courses?!

- a) We like to torture students. (We had to suffer through these courses, so all CS grads must suffer likewise.)
- b) Interviewers always ask OS questions, so you might as well know the answers.
- c) You will probably have to write OS code in the future.
- d) Understanding how the OS works is a fundamental concept in CS, and will help you become a better programmer/scientist.



Administrivia

- Instructor Contact:
 - Email: bogdan@cs.toronto.edu
 - Office: BA 4268 Office Hours: Tue 3:30-4:30, Thu 3-4:30
- Webpage:
 - <http://www.teach.cs.toronto.edu/~csc369h/fall>
- Lecture schedule:
 - Check it carefully – some lectures may be in tutorial time!
- Piazza:
 - Linked from course webpage. Read daily. Ask questions there first
 - Please come talk to me if you have any concerns about using Piazza
- Course Syllabus (due dates, policies, etc.):
 - Linked from course webpage
 - Must read it carefully!



Course prerequisites

- Make sure you have the prerequisites!
- If you don't have the prerequisites, ask me for a waiver by email (no guarantees though!)



Course Overview

- Four assignments writing code in C (37%)
 - A1: System calls (8%)
 - A2: Synchronization (8%)
 - A3: Virtual memory (9%)
 - A4: File systems (12%)
 - In pairs, but individual contributions will be checked via midterm and final
- Weekly tutorial exercises for marks (4%)
- Weekly in class exercises for marks (4%)
- Midterm (10%)
- Final exam (45%) – must get >40% to pass the course!

Late policy:
10 grace tokens
Each token = 2 hours



Weekly Exercises

- Strong evidence that people learn better or faster by doing rather than passively listening
- Exercises make it easier to connect lecture material (information delivery) to assignments and real world



Exercises

Exercises may have many forms:

- In class group exercises
- In class exercise, but option to hand in online
- Out of class exercise to be submitted online
- Most tutorials will have an exercise
- Tutorial exercises will be largely related to assignments

Marking Scheme:

- In class: Participation / effort based
- Tutorial: Correctness (and / or effort)
- Tutorial exercises will be best 8 out of 10
- In class exercises will be best 8 out of 10



Assignments

- Write good, professional code
- Comment it properly
- Debug it properly, find corner cases
- Solve problems as they come, find workarounds if needed
- Very important experience before getting a programming job
 - Please treat them as such!



Assignments

- Assignments are due at 10:00 p.m. on the due date - check website for final due dates
- Code **must work on teaching labs (formerly known as CDF)**
- **Make sure you commit all your source files; we cannot find files you never submitted**
- Code style matters!
- **Test-as-you-go**
- The code you submit has to work, even if it doesn't implement everything
- **Code that does not compile gets zero marks!**



Did you catch that?

I will not submit code that does not compile!
I will not submit code that does not compile!
I will not submit code that does not compile!
I will not submit code that does not compile!
I will not submit code that does not compile!
I will not submit code that does not compile!
I will not submit code that does not compile!
I will not submit code that does not compile!
I will not submit code that does not compile!
I will not submit code that does not compile!
I will not submit code that does not compile!





Assignments

- svn (make sure to revise, if necessary)
- Start early on the assignments!
 - Make sure you **can** commit in your repository; **commit often!**
 - **Do not wait** until the very last minute to submit your assignment!
 - **Read the submission instructions carefully** \Rightarrow penalties for incorrect submissions!
- **Must know your code for the entire assignment!**
 - For partners: **Work together**, even if you split the work!
 - May conduct **interviews** for some or all assignments!
 - **Not having a good understanding of all code** \Rightarrow 0 marks!





Start on assignments early!

- Strongly correlated with:
 - ✓ better grades
 - ✓ less stress
 - ✓ better understanding of the concepts
- VM setup (see tutorial)





Bird's eye view

	M	T	W	R	F	S	S	
Intro, Processes	11	12	13	14	15	16	17	
Syscalls, threads	18	19	20	21	22	23	24	→ A1 go-ahead
Concurrency, Semaphores, Locks	25	26	27	28	29	30	1	→ A1 due
Synchronization, CVs, monitors	2	3	4	5	6	7	8	→ A2 go-ahead
Scheduling	9	10	11	12	13	14	15	→ A2 due
Intro to VM, paging	16	17	18	19	20	21	22	
Midterm, Paging policies	23	24	25	26	27	28	29	→ A3 go-ahead
TLBs, Intro to File Systems	30	31	1	2	3	4	5	→ A3 due
No classes - Fall break	6	7	8	9	10	11	12	
File Systems	13	14	15	16	17	18	19	→ A4 go-ahead
FS recovery and fault tolerance	20	21	22	23	24	25	26	
Deadlocks	27	28	29	30	1	2	3	→ A4 due
Security, Exam review	4	5	6					



Holidays



Midterm



Don't Panic!

- Help is available in many forms
 - Lectures/tutorials: Ask questions!
 - Office hours: My time dedicated specifically to helping you
 - TA lab/office hours – get help with your code
 - Piazza: Faster response!
 - Email: Longer turnaround time
 - Anonymous email: for feedback
 - Undergraduate TA Help Center:

http://web.cs.toronto.edu/program/ugrad/ug_helpcentre.htm



Academic Integrity: Plagiarism

- Very serious academic offences, penalties are severe!
- Clear distinction between collaboration and cheating
 - Of course you can help your friend track down a bug. It is **never ok to submit code that is not your own, or to give someone else your code!**
 - Ask questions on Piazza, but don't add details about your solution (especially your code!)
- All potential plagiarism cases will be investigated fully
- We will run plagiarism detection software at the end of term!
 - Dropping the course does not get you out of trouble!
- Tips:
 - **Don't post your code in public places (Github, etc.)**
 - **Don't search for solutions - taking partial work or similar ideas will be detected!**
 - **Make sure your partner is not plagiarizing! Discuss your work together!**



Readings

Strongly Recommended!

Operating Systems: Three Easy Pieces

by Remzi H. Arpaci-Dusseau and
Andrea C. Arpaci-Dusseau

If you want more:

Modern Operating Systems

by Andrew Tannenbaum

Do the readings!



Introduction - Overview

- Introduction to this term's topics
- What is an OS, why it's important
- Recap some concepts you should know from 209 (e.g., processes)
- Why is it exciting to understand your OS, or work on designing OSes?
- Major OS components
 - Reason about approaches taken by OS designers



Introduction

- What is an OS and why do I want one?



Introduction

- What is an OS and why do I want one?
 - How does it relate to the other parts of a computer system?
 - Convenient abstraction of H/W
 - Protection, security, authentication
 - Communication
- Make sure to review some computer organization (258), systems concepts (209) and C concepts (209)!



Introduction

- What is an OS and why do I want one?
 - How does it relate to the other parts of a computer system?
 - Convenient abstraction of H/W
 - Protection, security, authentication
 - Communication
- *What are the major goals and components of an OS?*



Goals of the OS

- Primary: *convenience for the user*
 - It must be easier to compute with the OS than without it
- Secondary: *efficient* operation of the computer system
- The two goals are sometimes contradictory
 - Actually, often contradictory
 - Which goal takes precedence depends on the purpose of the computer system



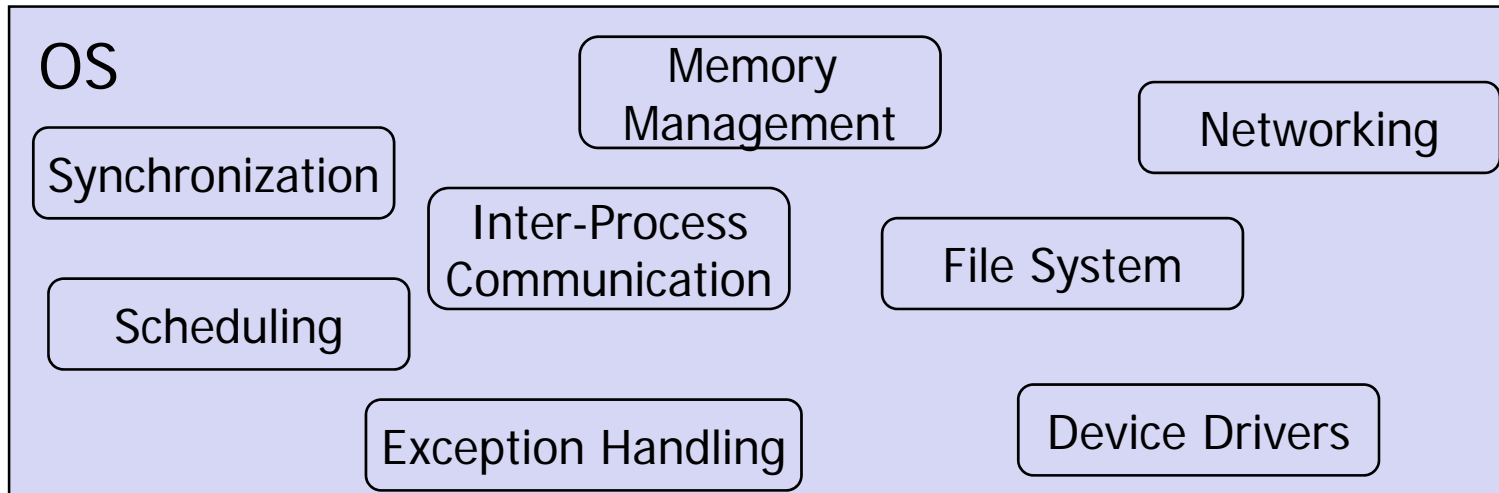
Roles of the OS

- An OS is a *virtual machine*
 - Extends and simplifies interface to physical machine
 - Provides a library of functions accessible through an API
- An OS is a *resource allocator*
 - allows the proper use of resources (hardware, software, data) in the operation of the computer system
 - provides an environment within which other programs can do useful work
- An OS is a *control program*
 - controls the execution of user programs to prevent errors and improper use of the computer
 - especially concerned with the operation and control of I/O devices

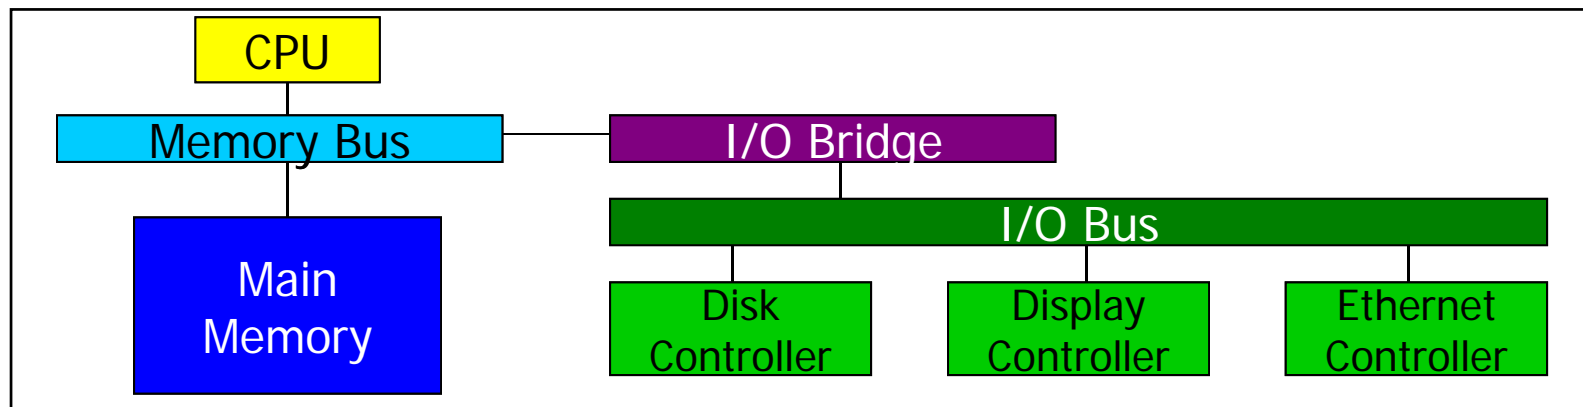


What is in the OS?

Software




Hardware





Storage Hierarchy

- processor registers, main memory, and auxiliary memory form a rudimentary memory hierarchy
- the hierarchy can be classified according to memory speed, cost, and volatility
- caches can be installed to hide performance differences when there is a large access-time gap between two levels





Major OS “Themes”

- Virtualization
 - Present physical resource as a more general, powerful, or easy-to-use form of itself
 - Present illusion of multiple (or unlimited) resources where only one (or a few) really exist
 - Examples: CPU, Memory (*demo*)
- Concurrency
 - Coordinate multiple activities to ensure correctness
- Persistence
 - Some data needs to survive crashes and power failures
- *Need abstractions, mechanisms, policies for all*



Next up...

- Hardware support for OS
- Bootstrapping
- Processes
 - What is a process?
 - Process lifecycle



Key Question: How to virtualize?

- Recall OS Goals:
 - Convenience for the user + efficient use of machine
- Virtualization helps with first goal
 - Virtual CPU + virtual memory gives each program the *illusion* that it owns the whole machine when it runs
- But how to do virtualization *efficiently*?
 - Consider Java language and Java Virtual Machine (JVM)
 - JVM provides *virtual instruction set* (Java bytecode)
 - JVM *interprets* each virtual instruction
 - *VERY* convenient, gain portability, not limited by HW
 - *VERY SLOW!*



OS Solution – Limited Direct Execution

- Key Abstraction: *the process*
 - Includes everything OS needs to know to manage running programs (more details later)
- Main Idea – Direct Execution
 - Set up CPU so that next instruction is fetched from code of process that should be executed
 - Let it go → no overhead during process execution
- Cool. Where does the “Limited” part come in?
 - Need to restrict operations process can perform
 - Need to regain control



Hardware Support for OSs

- Protection domains -> mode bit
- Interrupts
- Timers
- Memory Management unit
- Other hardware



Protection Domains

- Dual-mode operation: *user mode* and *system mode* (a.k.a *supervisor mode*, *monitor mode*, or *privileged mode*)
- Add a *mode bit* to the hardware and designate some instructions as *privileged instructions*
 - Intel actually has 4 “rings” for protection (CS register)
 - Attempting to execute privileged instruction from user mode causes *protection fault* → trap to OS
 - Protects the operating system from access by user programs, and protects user programs from each other
- So how can a process do something privileged? *System calls!*



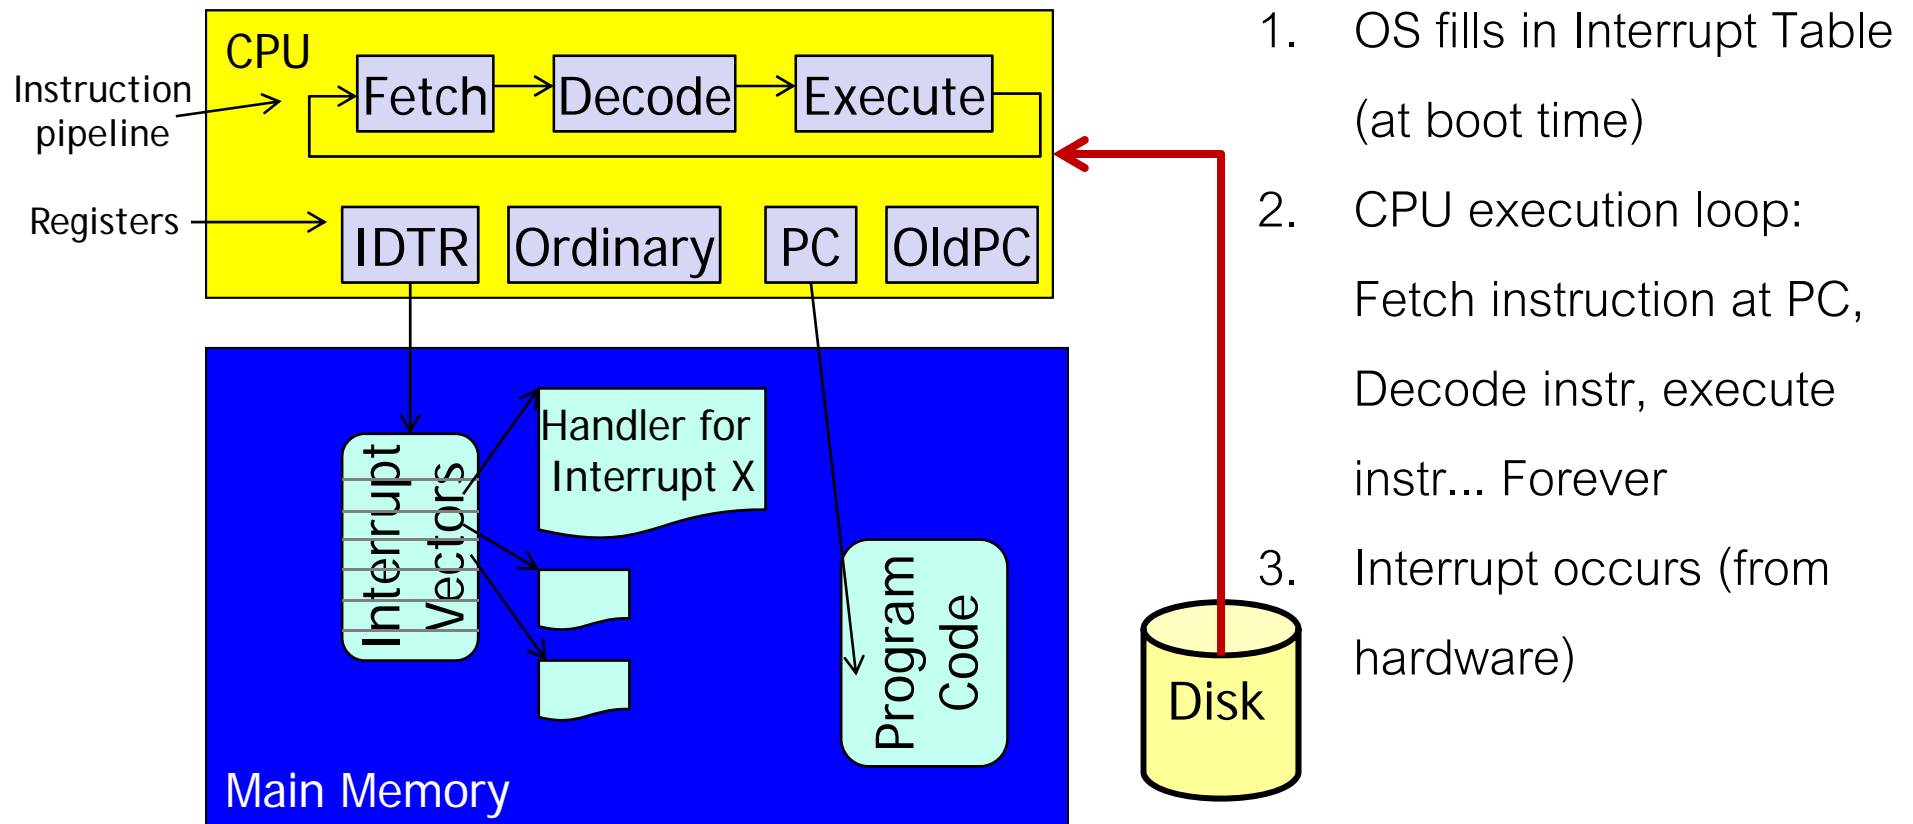
Protection Domains

- What instructions/operations would you expect to be privileged?
 - ✓ Setting mode bit?
 - ✓ Disabling interrupts?
 - ✓ Enabling interrupts?
 - ✓ Writing to device registers?
 - ✓ Performing DMA?
 - ✓ Halting the CPU?
- All of these are privileged!



Interrupts

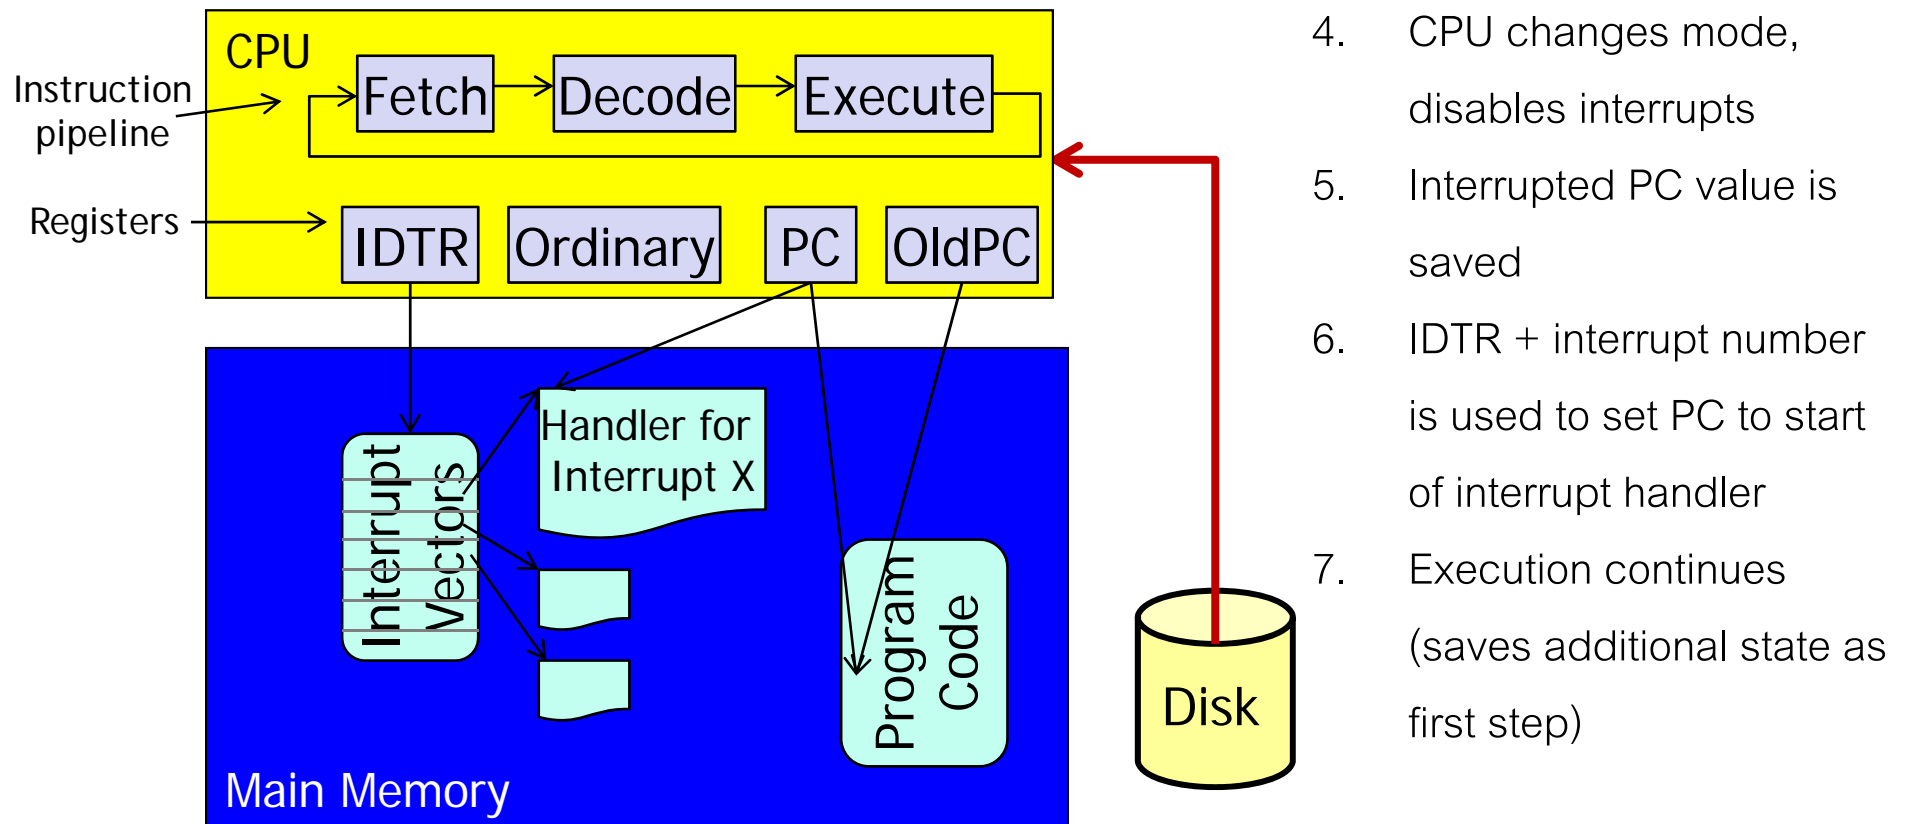
- Defn: a hardware signal that causes the CPU to jump to a pre-defined instruction called the *interrupt handler*
 - May be caused by software or hardware (Examples..)





Interrupts

- Defn: a hardware signal that causes the CPU to jump to a pre-defined instruction called the *interrupt handler*
 - May be caused by software or hardware (Examples..)





Interrupts and the OS

- Interrupt *mechanism* supports OS goal of *efficient virtualization* in two ways:
 1. Any illegal instruction executed by a process causes a software-generated interrupt (aka *an exception*)
 - OS gets control whenever the process does something it shouldn't
 - Attempting to execute a privileged instruction
 - Illegal memory access
 - Illegal instructions (e.g. divide by zero)
 2. Periodic hardware-generated timer interrupt ensures OS gets control back at regular intervals
 - Can switch processes to give virtual CPU illusion



Today

- Hardware support for OS
- Bootstrapping
- Processes



Bootstrapping

- Hardware stores small program in non-volatile memory
 - BIOS – Basic Input Output System
 - Knows how to access simple hardware devices
 - Disk, keyboard, display
- When power is first supplied, this program starts executing
- What does it do?



Operating System Startup

- Hardware starts in system mode (kernel mode), so OS code can execute immediately
- OS initialization:
 - Initialize internal data structures
 - Machine dependent operations are typically done first
 - Create first process (init)
 - Switch mode to user and start running first process
 - Wait for something to happen
 - “something” always starts with an interrupt
 - OS is entirely driven by external events
 - External to the OS, that is, not the entire computer



Today

- Hardware support for OS
- Bootstrapping
- Processes
 - Definition
 - Representation
 - Lifecycle
 - API



What is a Process?

- OS **abstraction for execution**
 - AKA a **job** or a **task** or a **sequential process**
- Definition: it's a **program in execution**
 - An active entity
- Programs are static entities with the ***potential*** for execution

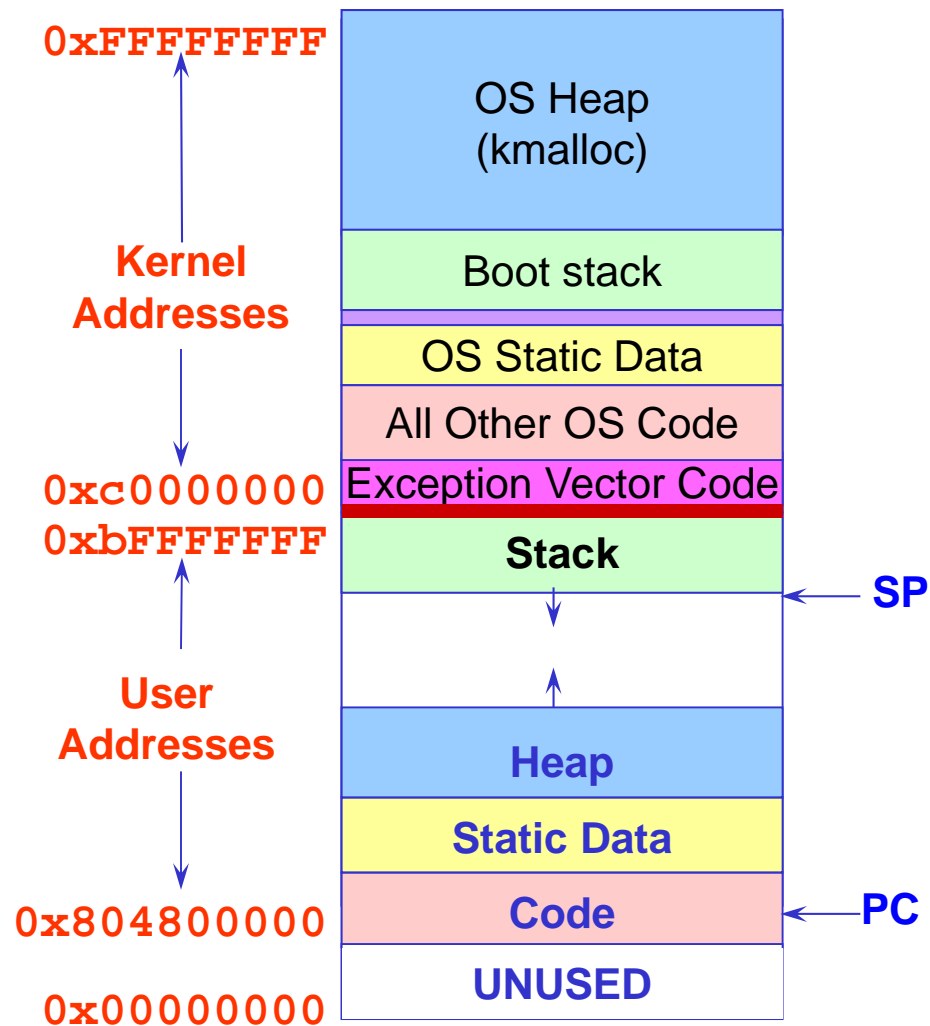


Program Layout in Memory

- What does a process's view of memory look like?
 - In other words, how is logical memory (logical address space) organized?
- BIG HINT: You saw this picture in CSC209 many times.



Program Layout in Memory





What is the OS view of a process?

- What data do we need to keep track of about a process?



OS View of a Process

- A *process* contains all of the state for a program in execution
 - An address space
 - The code and data for the executing user program + OS
 - An execution stack encapsulating the state of procedure calls
 - The program counter (PC) indicating the next instruction
 - A set of general-purpose registers with current values
 - A set of operating system resources
 - Open files, network connections, etc.
 - Context for kernel execution (a kernel thread & stack)
- A process is named using its process ID (**PID**)
- OS data about the process is stored in a *process control block (PCB)*



Process Control Block

- Generally includes:
 - process state (ready, running, blocked ...)
 - program counter: address of the next instruction
 - CPU registers: must be saved at an interrupt
 - CPU scheduling information: process priority
 - memory management info: page tables
 - accounting information: resource use info
 - I/O status information: list of open files



Linux PCB

- Called the *task_struct* in Linux
 - Defined in `/include/linux/sched.h`

```
struct task_struct {  
    /* these are hardcoded - don't touch */  
    volatile long state; /* -1 unrunnable, 0 runnable, >0 stopped */  
    long counter; long priority; unsigned long signal;  
    unsigned long blocked; /* bitmap of masked signals */  
    unsigned long flags; /* per process flags, defined below */  
    int errno; long debugreg[8]; /* Hardware debugging registers */  
    struct exec_domain *exec_domain;  
    /* various fields */  
    struct linux_binfmt *binfmt;  
    struct task_struct *next_task, *prev_task;  
    struct task_struct *next_run, *prev_run;  
    unsigned long saved_kernel_stack;  
    unsigned long kernel_stack_page;  
    int exit_code, exit_signal;  
    ...  
};
```



Keeping track of processes

How does the OS keep track of processes?

- Processes can be in various states (Ready, Running, Blocked)
- The OS maintains a collection of **state queues** that represent the state of all processes in the system
- Typically, the OS has one queue for each state
 - Ready, waiting for event X, etc.
- Each PCB is queued on a state queue according to its current state
- As a process changes state, its PCB is unlinked from one queue and linked into another



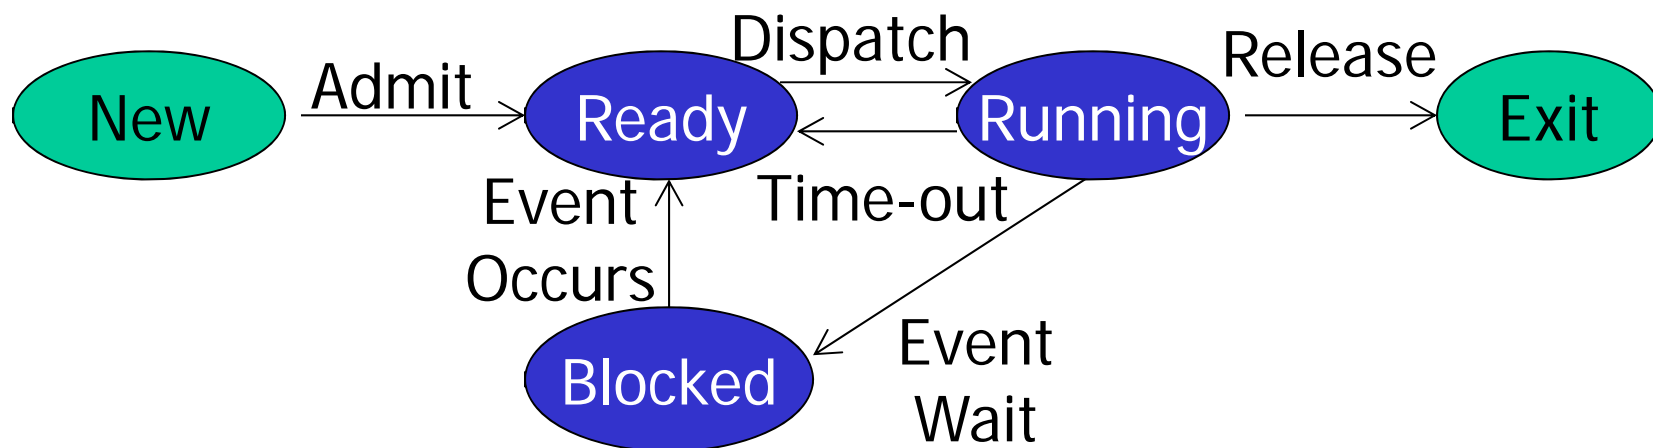
Today

- Processes
 - Definition
 - Representation
 - Lifecycle
 - API



Process states & state changes

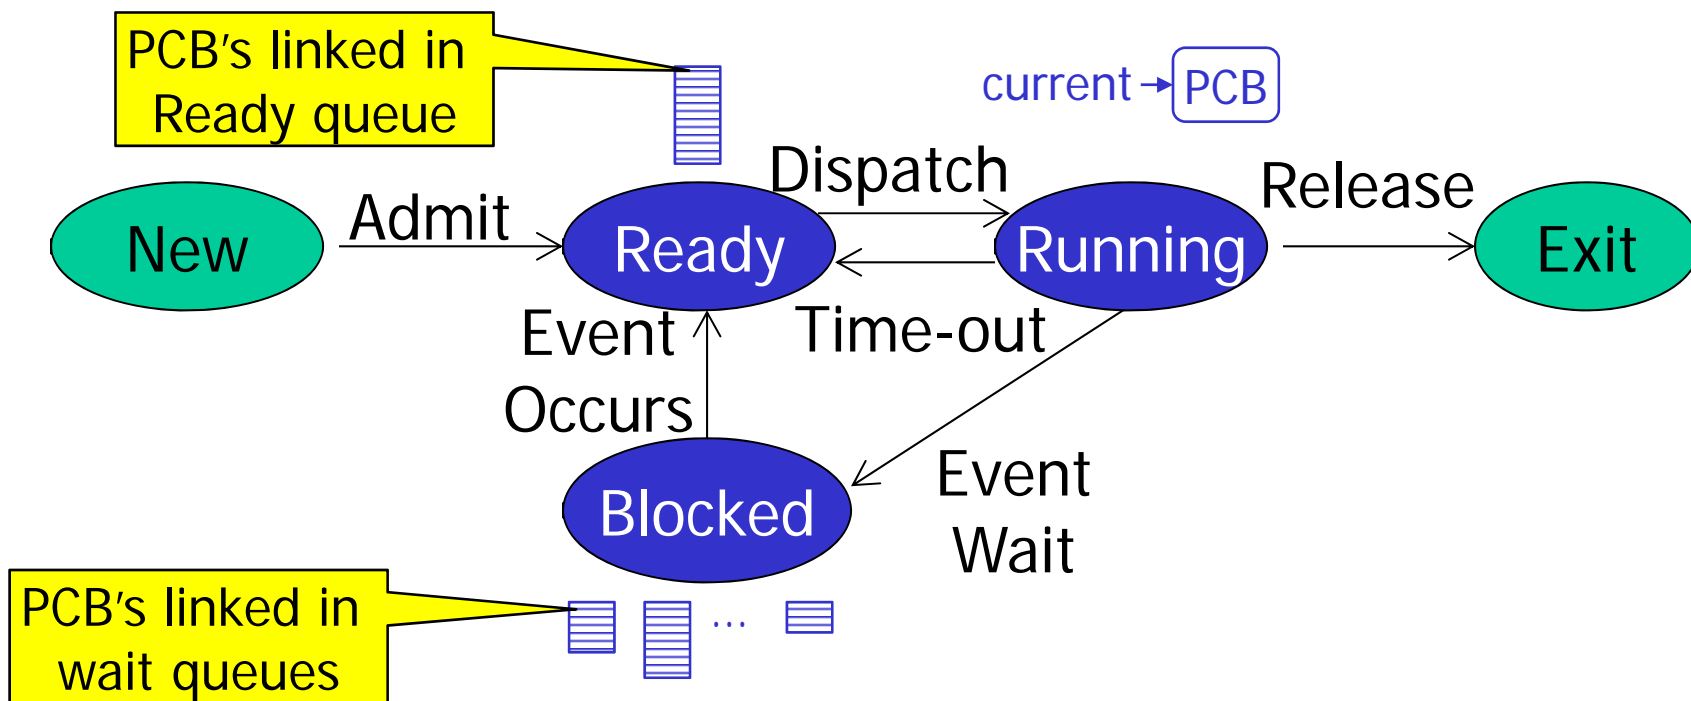
- The OS manages processes by keeping track of their *state*
 - Different *events* cause changes to a process state, which the OS must record/implement





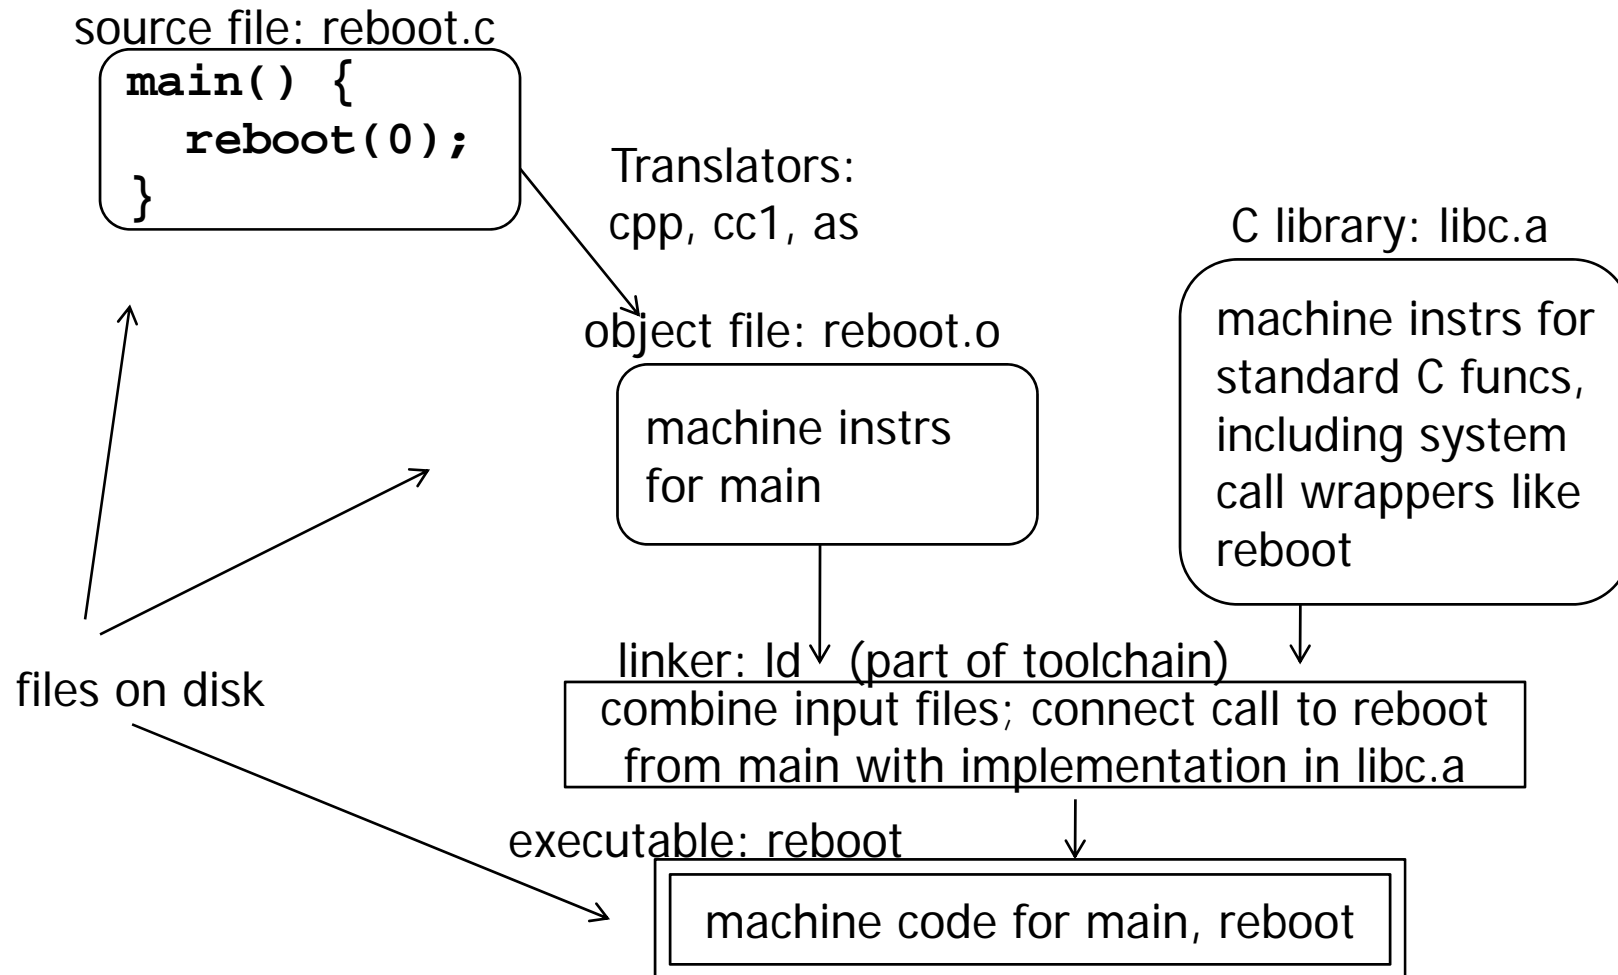
Process states & state changes

- The OS manages processes by keeping track of their *state*
 - Different *events* cause changes to a process state, which the OS must record/implement





From Program to Process... 1





From Program to Process... 2



-
1. Create new process
 - Create new PCB, user address space structure
 - Allocate memory
 2. Load executable
 - Initialize start state for process
 - Change state to “ready”
 3. Dispatch process
 - Change state to “running”



State Change: Ready to Running

- *context switch* == switch the CPU to another process by:
 - saving the state of the old process
 - loading the saved state for the new process
- When can this happen?
 - Process calls `yield()` system call (voluntarily)
 - Process makes other system call and is blocked
 - Timer interrupt handler decides to switch processes
 - Why would we ever need this?

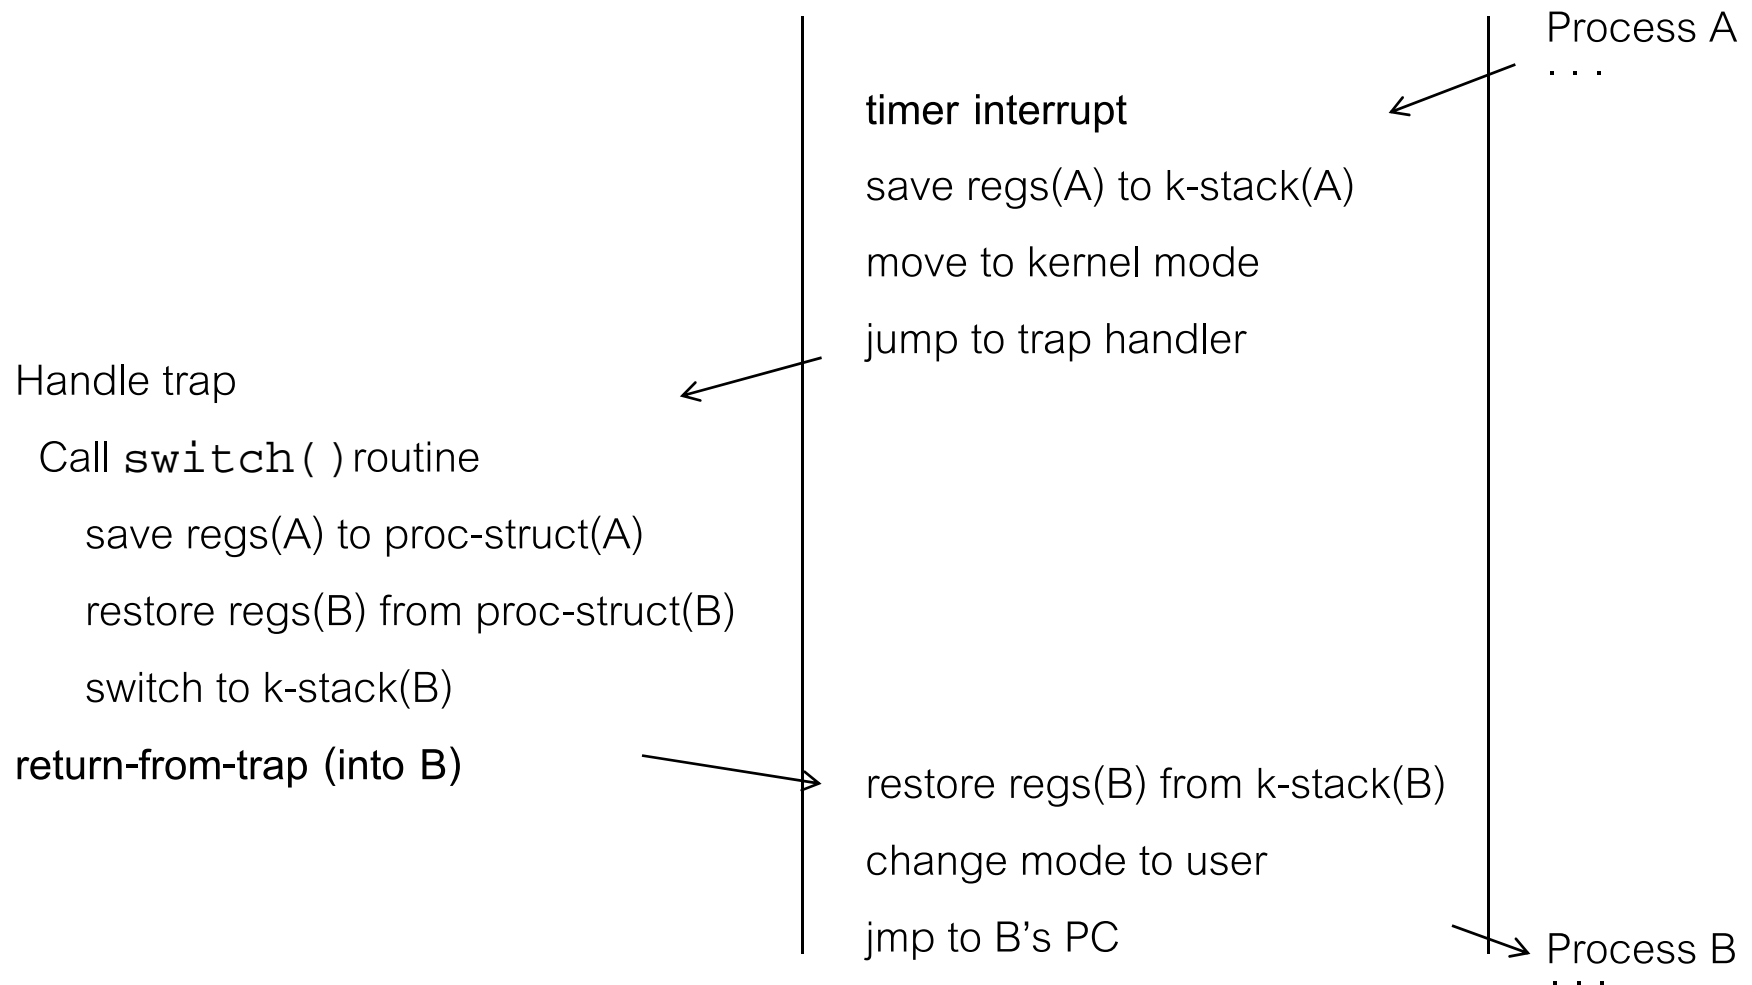


Context Switch (from Process A to B)

OS (kernel mode)

Hardware

User





Process Destruction

- `exit()`
- On `exit()`, a process voluntarily releases all resources
- But... OS can't discard everything immediately
- Why?
 - Must stop running the process to free everything
 - Requires *context switch* to another process
 - Parent may be waiting or asking for the return value
- `exit()` doesn't cause all data to be freed!
 - Zombies!



Zombies

- When a process exits, almost all of its resources are deallocated
 - Address space is freed, files are closed, etc.
- Some OS data structures retain the process's exit state
- The process retains its process ID (PID)
- It is a *zombie* until its parent cleans it up



*Source: Plants
vs Zombies*



Today

- Processes
 - Definition
 - Representation
 - Lifecycle
 - API



Process Creation: Unix

- In Unix, processes are created using `fork()` system call
 - `int fork()`
- `fork()`
 - Creates and initializes a new PCB
 - Creates a new address space
 - Initializes the address space with a **copy** of the entire contents of the address space of the parent
 - Initializes the kernel resources to point to the resources used by parent (e.g., open files)
 - Places the PCB on the ready queue
- Fork returns **twice**
 - Returns the child's PID to the parent, "0" to the child



Process Creation

- A process is created by another process
 - Parent is creator, child is created
 - In Linux, the parent is the “PPID” field of “ps -f”
 - Hierarchy: the first process (Unix): init (PID 1)
- In some systems, the parent defines (or donates) resources and privileges for its children
 - Unix: Process User ID is inherited – children of your shell execute with your privileges
- After creating a child, the parent may either wait for it to finish its task or continue in parallel
 - or continue for a while and then wait



Example: fork()

```
int main(int argc, char *argv[])
{
    char *name = argv[0];
    int child_pid = fork();
    if (child_pid == 0) {
        printf("Child of %s is %d\n", name, getpid());
        return 0;
    } else {
        printf("My child is %d\n", child_pid);
        return 0;
    }
}
```

What does this program print?



Process Creation: Unix (2)

- How do we actually start a new program?

```
int exec(char *prog, char *argv[])
```

- exec()
 - Stops the current process
 - Loads program “prog” into the process’ address space
 - Initializes hardware context and args for the new program
 - Places the PCB onto the ready queue
 - Note: It **does not** create a new process
- What does it mean for exec to return?



Inter-Process Communication (IPC)

- By design, processes are isolated from each other
 - But we often want them to exchange information
- OS provides various mechanisms for *inter-process communication*
 - Passing arguments to a newly exec'd program
 - Part of the `execv()` system call
 - Returning an integer exit status from child to parent
 - Part of the `waitpid()` and `exit()` system calls
 - Sending *signals* – a software analog of hardware interrupts
 - Provided by the `kill()` system call
 - Shared file system
 - Message passing, shared memory, synchronization primitives...