

CSC 369

Operating Systems

Lecture 3:

Synchronization: Semaphores & Locks




University of Toronto, Department of Computer Science



Assignment 1

- Remember: start early!





This Week

- Critical Section Problem
- Synchronization Primitives
 - Semaphores
 - Basic but flexible, Easy to understand
 - Can be hard to program with
 - Synchronization Hardware
 - Spinlocks
 - Very primitive, minimal semantics



Brief preview of scheduling

We have:

- Multiple threads/processes ready to run
- Some mechanism for switching between them
 - *Context switches*
- Some policy for choosing the next process to run
 - This policy may be *pre-emptive*
 - Meaning thread/process can't anticipate when it may be forced to yield the CPU
 - By design, it is not easy to detect it has happened (only the timing changes)



Synchronization

- Processes (and threads) interact in a multiprogrammed system
 - To share resources (such as shared data)
 - To coordinate their execution
- Arbitrary interleaving of thread executions can have unexpected consequences
 - We need a way to restrict the possible interleavings of executions
 - Scheduling is invisible to the application
- **Synchronization** is the mechanism that gives us this control



Flavours of Synchronization

- Two main uses:
 1. Enforce single use of a shared resource
 - Called the **critical section problem**
 - E.g. using a lock to ensure only one thread can print output to console at a time
 - T1: printf("Hello"); T2: printf("Goodbye");
 - Result should be "HelloGoodbye" OR "GoodbyeHello", but never "HeGooldbloye" or some other mixture
 2. Control order of thread execution
 - E.g. parent waits for child to finish
 - Ensure menu prints prompt *after* all output from thread running program



Motivating Example

- Suppose we write functions to handle withdrawals and deposits to bank account:

```
Withdraw(acct, amt) {  
    balance = get_balance(acct);  
    balance = balance - amt;  
    put_balance(acct, balance);  
    return balance;  
}
```

```
Deposit(acct, amt) {  
    balance = get_balance(acct);  
    balance = balance + amt;  
    put_balance(acct, balance);  
    return balance;  
}
```

- Now suppose you share this account with someone and the balance is \$1000
- You each go to separate ATM machines - you withdraw \$100 and your co-account holder deposits \$100



Example Continued

- We can represent this situation by creating separate threads for each action, which may run at the bank's central server:

```
Withdraw(acct, amt) {  
    balance = get_balance(acct);  
    balance = balance - amt;  
    put_balance(acct, balance);  
    return balance;  
}
```

```
Deposit(account, amount) {  
    balance = get_balance(acct);  
    balance = balance + amt;  
    put_balance(acct, balance);  
    return balance;  
}
```

- What's wrong with this implementation?
 - Think about potential schedules for these two threads



Interleaved Schedules

- The problem is that the execution of the two processes can be interleaved:

Schedule A

```
balance = get_balance(acct);  
balance = balance - amt;
```

```
balance = get_balance(acct);  
balance = balance + amt;  
put_balance(acct, balance);
```

```
put_balance(acct, balance);
```

Context
switch

Schedule B

```
balance = get_balance(acct);  
balance = balance - amt;
```

```
balance = get_balance(acct);  
balance = balance + amt;
```

```
put_balance(acct, balance);
```

```
put_balance(acct, balance);
```

- What is the account balance now?
- Is the bank happy with our implementation?
 - Are you?





What Went Wrong?

- Two concurrent threads manipulated a *shared resource* (the account) without any synchronization
 - Outcome depends on the order in which accesses take place
 - Aka a *race condition*
 - We need to ensure that only one thread at a time can manipulate the shared resource
 - So that we can reason about program behavior
- ➔ We need *synchronization*



Caution!

- Bank account problem can occur even with a simple shared variable, even on a uniprocessor:
 - T_1 and T_2 share variable X
 - T_1 increments X ($X := X+1$)
 - T_2 decrements X ($X := X-1$)
 - But at the machine level, we have:

T_1 :	LOAD X
	INCR
	STORE X

T_2 :	LOAD X
	DECR
	STORE X

- \Rightarrow Same problem of interleaving can occur!



Aside: What program data is shared?

That is, by threads in the same address space...

- Local variables are not shared (*private*)
 - Each thread has its own stack
 - Local vars are allocated on this private stack
 - Never pass/share/store a pointer to a local variable on another thread's stack!
- Global variables and static objects are *shared*
 - Stored in the static data segment, accessible by any thread
- Dynamic objects and other heap objs are *shared*
 - Allocated from heap with malloc/free (or new/delete, etc.)




Mutual Exclusion

- Given:
 - A set of n threads, T_0, T_1, \dots, T_{n-1}
 - A set of resources shared between threads
 - A segment of code which accesses the shared resources, called the *critical section, CS*
- We want to ensure that:
 - Only one thread at a time can execute in the critical section
 - All other threads are forced to wait on entry
 - When a thread leaves the CS, another can enter



The Critical Section Problem

- Design a protocol that threads can use to cooperate
 - Each thread must request permission to enter its CS, in its *entry* section
 - CS may be followed by an *exit* section
 - Remaining code is the *remainder* section



- Each thread is executing at non-zero speed
 - no assumptions about relative speed



Critical Section Requirements

1) Mutual Exclusion

- If one thread is in the CS, then no other is

2) Progress

- Only threads not in the “remainder” section can influence the choice of which thread enters next, and choice cannot be postponed indefinitely
- In other words: If no thread is in the CS, and some threads want to enter CS, they should be able to, unrestricted by threads in the “remainder”

3) Bounded waiting (no starvation)

- If some thread T is waiting on the CS, then there is a limit on the number of times other threads can enter CS before this thread is granted access

- **Performance**

- The overhead of entering and exiting the CS is small with respect to the work being done within it



Some Assumptions & Notation

- Assume **no special hardware instructions** (no H/W support)
- Assume no restrictions on the # of processors (for now)
- Assume that basic machine language instructions (LOAD, STORE, etc.) are **atomic**:
 - If two such instructions are executed concurrently, the result is equivalent to their sequential execution in some unknown order
 - On modern architectures, this assumption may be false
- Let's consider a simple scenario: only 2 threads, numbered T_0 and T_1
 - Use T_i to refer to one thread, T_j for the other ($j=1-i$) when the exact numbering doesn't matter
- Let's look at one solution... [Exercise]



2-Thread Solutions: Take 1

- Let the threads share an integer variable *turn* initialized to 0 (or 1)
- If $turn=i$, thread T_i is allowed into its CS

```
My_work(id_t id) { /* id_t can be 0 or 1 */  
    ...  
    while (turn != id) ; /* entry section */  
    /* critical section, access protected resource */  
    turn = 1 - id;      /* exit section */  
    ...                /* remainder section */  
}
```



Only one thread at a time can be in its CS



Progress is not satisfied

- Requires strict alternation of threads in their CS: if $turn=0$, T_1 may not enter, even if T_0 is in the remainder section



2-Thread Solutions: Take 2

- First attempt does not have enough info about state of each process. It only remembers which process is allowed to enter its CS
- Replace turn with a shared flag for each thread
 - **boolean** `flag[2]` = {false, false}
 - Each thread may update its own flag, and read the other thread's flag
 - If `flag[i]` is true, T_i is ready to enter its CS
- Exercise ..



A Closer Look at 2nd Attempt

```
My_work(id_t id) { /* id can be 0 or 1 */
    ...
    while (flag[1-id]) ; /* entry section */
    flag[id] = true;      /* indicate entering CS */
    /* critical section, access protected resource */
    flag[id] = false;     /* exit section */
    ...                  /* remainder section */
}
```

- Progress guaranteed?
- Starvation?
- Mutual exclusion is not guaranteed
 - Each thread executes *while* statement, finds *flag* set to false
 - Each thread sets own *flag* to *true* and enters CS



Example Execution Sequence

- Thread0 (id = 0)

- Thread1 (id = 1)

while (flag[1]);
/*false*/

switch

while (flag[0]);
/*false*/

flag[id] = true;
/* in crit. sect. */

flag[id] = true;
/* in crit. sect. */

flag[id] = false;

flag[id] = false;

Time
↓

Can't fix this by changing order of testing and setting *flag* variables
(leads to *deadlock*)



2-Thread Solutions: Take 3

- Combine key ideas of first two attempts for a correct solution
- The threads share the variables *turn* and *flag* (where *flag* is an array, as before)
- Basic idea:
 - Set own flag (indicate interest) and set turn to self
 - Spin waiting while turn is self AND other has flag set (is interested)
 - If both threads try to enter their CS at the same time, *turn* will be set to both 0 and 1 at roughly the same time. Only one of these assignments will last. The final value of *turn* decides which of the two threads is allowed to enter its CS first.
- This is the basis of *Dekker's Algorithm* (1965) and *Peterson's Algorithm* (1981) - Modern OS book, 4th Ed. (A. Tanenbaum) – Fig 2.24, p 125



Peterson's Algorithm

```
int turn;
int flags[2];  /* Shows "interest" in the CS,
                Are both initially 0, aka false. */

My_work(id_t id) {  /* id can be 0 or 1 */
    ...
    flag[id] = true;
    turn = id;
    while (turn == id && flag[1-id]) ;
    /* critical section, access protected resource */
    flag[id] = false;  /* exit section */
    ...                /* remainder section */
}
```

- Convince yourself that this works...




Multiple-Thread Solutions

- Peterson's Algorithm can be extended to N threads
- Another approach is Lamport's *Bakery Algorithm*
 - Upon entering each customer (thread) gets a #
 - The customer with the lowest number is served next
 - No guarantee that 2 threads do not get same #
 - In case of a tie, thread with the lowest id is served first
 - Thread ids are unique and totally ordered
 - Mutual exclusion? Progress guaranteed? Starvation?



How multithreaded programming feels like ...

Multithreaded programming



Source: 9gag.com



Next: Higher-level Abstractions for CS's

- Semaphores
 - Basic, easy to understand, hard to program with
- Locks
 - Very primitive, minimal semantics
- Condition variables
 - Stronger semantics, easier for diverse conditions
- Monitors
 - High-level, ideally has language support (Java)




Semaphores

- Semaphores are abstract data types that provide synchronization.

- They include:

- An integer **counter** variable, accessed only through 2 atomic operations
- The atomic operation **wait** (also called ***P*** or ***decrement***) - decrement the variable and block until semaphore is free
- The atomic operation **signal** (also called ***V*** or ***increment***) - increment the variable, unblock a waiting thread if there are any
- A queue of **waiting** threads



Source: wikipedia.org



Types of Semaphores

- Mutex (or Binary) Semaphore (count = 0/1)

- Single access to a resource
- Mutual exclusion to a critical section



One! One thread
in the critical
section,
Ah ah ah!

Source: muppethub.com

- Counting semaphore

- A resource with many units available, or a resource that allows certain kinds of unsynchronized concurrent access (e.g., reading)
- Multiple threads can pass the semaphore
- Max number of threads is determined by semaphore's initial value, *count*
 - Mutex has *count* = 1, counting has *count* = *N*



Using Binary Semaphores

Have semaphore, S, associated with acct

- Consider its initial value is 1

```
typedef struct account {  
    double balance;  
    semaphore S;  
} account_t;  
  
Withdraw(account_t *acct, amt) {  
    double bal;  
    wait(acct->S);  
    bal = acct->balance;  
    bal = bal - amt;  
    acct->balance = bal;  
    signal(acct->S);  
    return bal;  
}
```

Three threads execute Withdraw()

```
wait(acct->S);  
bal = acct->balance;  
bal = bal - amt;
```

```
wait(acct->S);
```

```
wait(acct->S);
```

```
acct->balance = bal;  
signal(acct->S);
```

```
...  
signal(acct->S);
```

```
...  
signal(acct->S);
```

It is undefined which thread runs after a signal



Atomicity of `wait()` and `signal()`

- We must ensure that two threads cannot execute *wait* and *signal* at the same time
- This is another critical section problem!
 - Use lower-level primitives to implement wait and signal
 - Uniprocessor: disable interrupts
 - Multiprocessor: use hardware instructions
 - Examples later in this lecture...



Locks vs. Semaphores

- A binary semaphore (with initial value 1) can be used just like a lock
- Why bother with both abstractions?
 - Semantic difference – logically, a lock has an “owner” and can only be released by its owner
 - Permits some error checking
 - Helps reason about the correct behavior
- Let's look at a synchronization problem...



Producer / Consumer

- Classic synchronization problem
- Think how you would implement a pipe

```
Producer() {  
    while(1) {  
  
        write();  
  
    }  
}
```

```
Consumer() {  
    while(1) {  
  
        read();  
  
    }  
}
```

- To be continued later ..




Posix Semaphores API

```
#include <semaphore.h>
sem_t mysem;
sem_init(&mysem, 0, VALUE);
sem_wait(&mysem);    → s.wait or P()
sem_post(&mysem);    → s.signal or V()
```




Example: semaphores and locks (pthread_mutex)





And now ...

- Semaphore exercise (use wait/signal)...

s1(0) s2(0)

Thread A

a1

s1.signal()

s2.wait()

a2

Thread B

b1

s1.wait()

s2.signal()

b2



Synchronization Hardware

- To build these higher-level abstractions, it is useful to have some help from the hardware
- On a uniprocessor, in the OS:
 - disable interrupts before entering critical section (prevents context switches)

```
oldspl = splhigh();  
CRITICAL SECTION CODE  
splx(oldspl);
```

- Disabling interrupts is insufficient on a multiprocessor. Why?
- Need some special atomic instructions



Atomic Instructions: Test-and-Set

- The semantics of test-and-set are:
 - Record the old value of the variable
 - Set the variable to some non-zero value
 - Return the old value
- Hardware executes this atomically!
- Can be used to implement simple lock variables

```
boolean test_and_set(boolean *lock) {  
    boolean old = *lock;  
    *lock = True;  
    return old;  
}
```



Alternate Defn of Test-and-Set

- We'll use "TAS" in the code example for "test-and-set"

```
boolean TAS(boolean *lock) {  
    boolean old = *lock;  
    *lock = True;  
    return old;  
}
```

```
boolean TAS(boolean *lock) {  
    if(*lock == False) {  
        *lock = True;  
        return False;  
    } else  
        return True;  
}
```

- *lock* is always *True* on exit from test-and-set
 - Either it was *True* (locked) already, and nothing changed, or it was *False* (available), but the caller now holds it
- Return value is either *True* if it was locked already, or *False* if it was previously available



A Lock Implementation

- There are two operations on locks: *acquire()* and *release()*

```
boolean lock;  
  
void acquire(boolean *lock) {  
    while(test_and_set(lock));  
}  
  
void release(boolean *lock) {  
    *lock = false;  
}
```

When false, we know
that we've acquired it

To release, simply
turn it to false.

- This is a *spinlock*
 - Uses *busy waiting* - thread continually executes *while* loop in *acquire()* , consumes CPU cycles



Using Locks

Function Definitions

```
Withdraw(acct, amt) {  
    acquire(lock);  
    balance = get_balance(acct);  
    balance = balance - amt;  
    put_balance(acct, balance);  
    release(lock);  
    return balance;  
}
```

```
Deposit(account, amount) {  
    acquire(lock);  
    balance = get_balance(acct);  
    balance = balance + amt;  
    put_balance(acct, balance);  
    release(lock);  
    return balance;  
}
```

Possible schedule

```
acquire(lock);  
balance = get_balance(acct);  
balance = balance - amt;
```

```
acquire(lock);
```

```
put_balance(acct, balance);  
release(lock);
```

```
balance = get_balance(acct);  
balance = balance + amt;  
put_balance(acct, balance);  
release(lock);
```



More Special Instructions

- *Swap* (or *Exchange*) instruction
 - Operates on two words atomically
 - Can also be used to solve critical section problem
 - `Swap(boolean *varA, boolean *varB)`

```
boolean lock = false; // shared by all processes  
...
```

```
// in each thread  
boolean key = true; // local to each thread  
while(key)  Swap(&lock, &key); // ENTRY  
  
// Critical section  
  
Swap(&lock, &key); // EXIT
```




Other considerations

- Spinlocks are built on machine instructions
- Machine instructions have three problems:
 - Busy waiting
 - Starvation is possible
 - when a thread leaves its CS, the next one to enter depends on scheduling
 - a waiting thread could be denied entry indefinitely
 - Deadlock is possible through *priority inversion*
 - *More on that later, with scheduling...*



Sleep Locks

- Instead of spinning, put thread to sleep (into “blocked” state) while waiting to acquire a lock
- Requires a queue for waiting threads
 - Linux: wait queues



Source: <http://joyreactor.com>

```
wait_event(queue, condition)
wake_up(wait_queue_head_t *queue);
```