

CSC 369

Week 8: Paging Design & Implementation



University of Toronto, Department of Computer Science



Agenda

- Recap of full address translation
- Other implementation concerns
- Advanced Functionality



Thrashing

- Page replacement algorithms should avoid **thrashing**
 - When more time is spent by the OS in paging data back and forth from disk than executing user programs
 - No time spent doing useful work (making progress)
 - In this situation, the system is **overcommitted/oversubscribed**
 - No idea which pages should be in memory to reduce faults
 - Could just be that there isn't enough physical memory for all of the processes in the system
 - Ex: Running Windows10 with 4MB of memory...
 - Possible solutions
 - Swapping – write out all pages of a process and suspend it
 - OOM Killer daemon
 - Buy more memory



CPU utilization

- What percentage of time is the CPU busy
 - Mostly I/O heavy processes – CPU utilization is low
 - Compute-intensive processes – CPU utilization is high
- Say CPU utilization is low. Is it a good idea to increase the degree of multiprogramming, to better utilize the idle CPU cycles?
 - Nope! In fact typically decreases CPU utilization
 - Less memory is available to each program => higher page fault likelihood
- What about decreasing?



Page Buffering

- Previously, we assumed the replacement algorithm is run and a victim page selected *when a new page needs to be brought in*
- Most of these algorithms are too costly to run on every page fault
 - Maintain a pool of free pages (free page list)
 - Run replacement algorithm when pool becomes too small (“low water mark”), free enough pages to at once replenish pool (“high water mark”)
 - Uses dedicated kernel thread, *the paging daemon*
 - On page fault, grab a frame from the free list
 - Frames on free list still hold previous contents, can be “rescued” if virtual page is referenced before reallocation



Addressing Page Tables

Where do we store page tables (which address space)?

- Physical memory
 - Easy to address, no translation required (or very simple translation, like linear offset)
 - allocated page tables consume memory for lifetime of Virt. AS
- Virtual memory (OS virtual address space, KSEG2)
 - Cold (unused) page table pages can be paged out to disk
 - But, addressing page tables requires translation
 - How do we stop recursion?
 - Do not page the outer page table (called **wiring**)
- If we're going to page the page tables, might as well page the entire OS address space, too
 - Need to wire special code and data (fault, interrupt handlers)



Managing Swap Space

- Option 1: Use raw disk partition
- Option 2: Use ordinary large file in file system
- Tradeoffs?
- When should swap be allocated / freed?
 - On process startup / shutdown?
 - On pageout / pagein?



Address Translation Redux

- We started this topic with the high-level problem of translating virtual addresses into physical address
- We've covered all of the pieces
 - Virtual and physical addresses
 - Virtual pages and physical page frames
 - Page tables and page table entries (PTEs), protection
 - TLBs
 - Demand paging
- Now let's put it together, bottom to top



The Common Case

- Situation: Process is executing on the CPU, and it issues a **read** to an address
 - What does this address typically contain?
 - What kind of address is it? Virtual or physical?



Read Access (Load)

The read address goes to the TLB in the MMU

1. TLB does a lookup using the **page number** of the address
 2. Common case is that the page number matches, returning a **page table entry (PTE)** for the mapping for this address
 3. TLB validates that the **PTE protection** allows reads
 4. PTE specifies which **physical frame** holds the page
 5. MMU combines physical frame & offset into a **physical address**
 6. MMU reads from that physical addr, returns value to CPU
- Note: **This is all done by the hardware**



TLB Misses

- At steps 2 or 3, two other things can happen
 1. TLB does not have a PTE mapping for this virtual address
 2. PTE exists, but memory access violates PTE protection bits
- We'll consider each in turn



1. TLB does not have mapping

- Two possibilities:
 1. MMU loads PTE from page table in memory
 - Hardware managed TLB, OS not involved in this step
 - OS has already set up the page tables so that the hardware can access it directly
 2. Trap to OS
 - Software managed TLB
 - OS does lookup in page table, loads PTE into TLB
 - Return from exception, retry memory access
 - Note: Most machines will only support one method or the other
- Now there is a PTE for the address in the TLB
 - Known as a *minor page fault* (no I/O needed)



2. Access not permitted by PTE

- PTE can indicate a **protection fault**
 - **Read/write/execute** – operation not permitted on page
 - **Invalid** – virtual page not allocated, or page not in physical memory
- TLB traps to the OS (software takes over)
 - **R/W/E** – OS may send fault back up to process, or might be using protection for other purposes (e.g., copy on write, mapped files)
 - **Invalid**
 - Virtual page not allocated in address space
 - OS sends fault to process (e.g., segmentation fault)
 - Page not in physical memory (known as *major page fault*)
 - OS allocates frame and reads it in



Check page faults

- Example, on CDF:
 - Number of major and minor page faults on a Linux system (all processes!)
 - `ps -eo min_flt,maj_flt,cmd`
 - Check the page faults generated by a specific program during its execution, e.g.:
 - `/usr/bin/time -v firefox`
- when program terminates, see the stats (Major/Minor page faults)



Next up

- Other implementation concerns
 - How much memory should each process get?
- Advanced Functionality
 - Shared memory
 - Copy-on-write
 - Memory mapped files




Fixed vs. Variable Space Allocation

- In a multiprogramming system, we need a way to allocate memory to competing processes
- **Problem:** How to determine how much memory to give to each process?
 - *Fixed space algorithms*
 - Each process is given a limit of pages it can use
 - When it reaches the limit, it replaces from its own pages
 - **Local replacement**
 - Some processes may do well while others suffer
 - *Variable space algorithms*
 - Process' set of pages grows and shrinks dynamically
 - **Global replacement** - one process can ruin it for the rest
 - Local replacement – replacement, set size are separate for each process



Working Set Model

- A **working set** of a process is used to model the dynamic locality of its memory usage
 - Defined by Peter Denning in 60s
- Definition
 - $WS(t, \Delta) = \{\text{pages } P \text{ such that } P \text{ was referenced in the time interval } (t, t - \Delta)\}$
 - $t = \text{time}$, $\Delta = \text{working set window (measured in page refs)}$
- A page is **in the working set (WS) only if it was referenced in the last Δ references**





Working Set Size

- The **working set size** is the number of pages in the working set
 - The number of pages referenced in the interval $(t, t-\Delta)$
- The working set size **changes with program locality**
 - During periods of poor locality, you reference more pages
 - Within that period of time, the working set size is larger
- Intuitively, want the working set to be the set of pages a process needs **in memory to prevent heavy faulting**
 - Each process has a parameter Δ that determines a working set with few faults
 - Denning: Don't run a process unless working set is in memory



Working Set Problems

- Problems
 - How do we determine Δ ?
 - How do we know when the working set changes?
- Too hard to answer
 - So, working set is not used in practice as a page replacement algorithm
- However, it is still used as an abstraction
 - The intuition is still valid
 - When people ask, “How much memory does Firefox need?”, they are in effect asking for the size of Firefox’s working set
- Approximations may be useful in practice



Advanced Functionality

- Some advanced functionality that the OS can provide applications using virtual memory tricks
 - Shared memory
 - Copy on Write




Sharing

- Private virtual address spaces protect applications from each other
 - Usually exactly what we want
- But this makes it difficult to share data (have to copy)
 - Parents and children in a forking Web server or proxy will want to share an in-memory cache without copying
- We can use **shared memory** to allow processes to share data using **direct memory references**
 - Both processes see updates to the shared memory segment
 - Process B can immediately read an update by process A
 - How are we going to **coordinate access to shared data**?



Sharing (2)


- How can we implement sharing using page tables?
 - Have PTEs in both tables map to the same physical frame
 - Each PTE can have different protection values
 - Must update both PTEs when page becomes invalid





Mapping Shared Regions

- Can map shared memory at **same** or **different** virtual addresses in each process' address space
 - **Different**: Flexible (no address space conflicts), but pointers inside the shared memory segment are invalid (Why?)
 - **Same**: Less flexible, but shared pointers are valid (Why?)





Copy on Write

- OSes spend a lot of time copying data
 - System call arguments between user/kernel space
 - Entire address spaces to implement `fork()`
- Use **Copy on Write (CoW)** to defer large copies as long as possible, hoping to avoid them altogether
 - Instead of copying pages, create **shared mappings** of parent pages in child virtual address space
 - Shared pages are protected as read-only in child
 - Reads happen as usual
 - Writes generate a protection fault, trap to OS, copy page, change page mapping in client page table, restart write instruction
 - How does this help `fork()`? (Implemented as Unix `vfork()`)



Memory mapped files

- Mapped files enable processes to do file I/O using loads and stores
 - Instead of “open file, read into buffer, operate on buffer, ...”
- Bind a file to a virtual memory region (`mmap()` in Unix)
 - PTEs map virtual addresses to physical frames holding file data
 - Virtual address base + N refers to offset N in file
 - => Can read or write at various offsets in file, using memory operations
- Initially, all pages mapped to file are invalid
 - OS reads a page from file when invalid page is accessed
 - OS writes a page to file when evicted, or region unmapped



Memory mapped files

- File is essentially backing store for that region of the virtual address space (instead of using the swap file)
 - Virtual address space not backed by “real” files also called “Anonymous VM”
- Advantages
 - Uniform access for files and memory (just use pointers)
 - Less copying
- Drawbacks
 - Process has less control over data movement
 - OS handles faults transparently
 - Does not generalize to streamed I/O (pipes, sockets, etc.)



Summary

Paging mechanisms:

- Optimizations
 - Managing page tables (space)
 - Efficient translations (TLBs) (time)
 - Demand paged virtual memory (space)
- Recap address translation

Paging policies

Advanced Functionality

- Sharing memory
- Copy on Write
- Memory mapped files



Example Paging Systems

- Next few slides cover Windows and Linux VM *very* briefly
- For fun/extra knowledge only – not on exam



Windows XP Virtual Memory

- 4KB page size on IA32 processors
 - 8 kB on the IA64
- 4GB virtual address space, upper 2 GB used by XP in kernel mode
- Multi-level page table
 - Page directory contains 1024 page directory entries (PDE) of size 4 bytes
 - PDEs point to page tables containing 1024 page table entries (PTEs) of size 4 bytes
- Page frames are tracked using a “page frame database” with one entry per page of physical memory; entry points to PTE which points to frame



Windows XP Paging Policy

- Local page replacement
 - Per-process FIFO
 - Pages are stolen from processes using more than their minimum working set
 - Processes start with a default of 50 pages
 - XP monitors page fault rate and adjusts working-set size accordingly
 - On page fault, *cluster* of pages around the missing page are brought into memory



Linux Paging

- Global replacement, like most Unix
- Modified second-chance clock algorithm
 - Pages *age* with each pass of the clock hand
 - Pages that are not used for a long time will eventually have a value of zero
- Continually under development



Some A3 tips

- Main goal: implement 2-level page tables and demand paging
- We use a simulator to simulate the physical memory
 - 1. We generate memory access traces of some programs (like in Exercise 8)
 - 2. We go through the trace of virtual addresses, translate each of them into a physical address, using 2-level page table and demand paging (swapping)
 - 3. Implement different page replacement policies
 - 4. Analyze memory access patterns



More A3 tips

- 32 vs 64-bit traces (TRACE_64 macro)

28 bits unused	12 bits PGDIR_INDEX (4K PDEs)	12 bits PGTBL_INDEX (4K PTEs per tbl)	12 bits OFFSET (4K page size)
----------------	-------------------------------------	---	-------------------------------------

64-bit

$\text{PGDIR_SHIFT} = 24 \text{ bit} \Rightarrow$
 $\text{PGDIR_INDEX} = \text{VADDR} \gg 24$

10 bits PGDIR_INDEX (1K PDEs)	10 bits PGTBL_INDEX (1K PTEs per tbl)	12 bits OFFSET (4K page size)
-------------------------------------	---	-------------------------------------

32-bit

$\text{PGDIR_SHIFT} = 22 \text{ bit} \Rightarrow$
 $\text{PGDIR_INDEX} = \text{VADDR} \gg 22$

Your code should be the same regardless of 32/64-bit. Do not use hard-coded numbers, use the macros provided in “sim.h”!