

CSC 369

Week 11:

Example File Systems

Reliability and Write Optimizations



University of Toronto, Department of Computer Science



Overview

- Last time:
 - Optimizations: caching, read-ahead
 - Disk characteristics, optimizations
 - Disk-aware allocation, I/O scheduling
- This week:
 - Example file systems
 - Crash consistency and recovery
 - Optimizing writes
 - Log-structured file systems
 - VFS
 - Solid State Drives (SSDs)

Example File Systems



University of Toronto, Department of Computer Science



FS Comparison

	FAT	FFS	NTFS
Index structure	Linked list	Tree (inodes)	Tree (extents)
Index structure granularity	Block	Block	Extent
Free space management	FAT array	Bitmap	Bitmap
Locality heuristics	Defragmentation	Block groups, reserved space	Best fit Defragmentation

Table Source: Operating Systems, Anderson & Dahlin, p 554

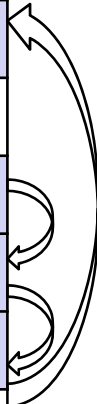


FAT

- File Allocation Table
 - Late 70s
 - Blocks allocated in a linked index structure

FAT

0	
1	
2	
3	nil
4	
5	6
6	7
7	3



Directory

afile	5



FAT

- Directories map file name to first block of file
- FAT stores both the linked list of blocks belonging to files and the free blocks
- Limitations
 - Poor random access
 - Poor locality
 - Limited file metadata and access control



Ext2, Ext3, Ext4

- Linux file system evolution
- Ext2 originally borrowed heavily from FFS
- Recall: Reduce seeks for faster reads, etc.
- More details on reliability and optimizations for writes




Reliability and Write Optimizations

- How do we guarantee consistency of on-disk storage?
- How do we handle OS crashes and disk errors?
- How do we optimize writes?



FFS: Consistency Issues - Overview

- Inodes: fixed size structure stored in cylinder groups



- Metadata updates must be synchronous operations. Why?
- File system operations affect multiple metadata blocks
 - Write newly allocated inode to disk before its name is entered in a directory.
 - Remove a directory name before the inode is deallocated
 - Deallocate an inode (mark as free in bitmap) before that file's data blocks are placed into the cylinder group free list.




FFS Observation 1: Crash recovery

- If the OS crashes in between any of these synchronous operations, then the file system is in an inconsistent state.
- Solutions (overview):
 - *fsck* – post-crash recovery process to scan file system structure and restore consistency
 - All data blocks pointed to by inodes (and indirect blocks) must be marked allocated in the data bitmap
 - All allocated inodes must be in some dir entries
 - Inode link count must match
 - Log updates to enable roll-back or roll-forward




Example: update

- Consider a simple update: append 1 data block to a file
- Assume a similar FS structure as seen before:



- Add a data block: Db .. What changes?



- Three writes: I[v2], B[v2], Db



Crash consistency

- What if only one write succeeds before a crash?
 - 1. **Just Db write succeeds.**
 - No inode, no bitmap => as if the write did not occur
 - FS not inconsistent, but data is lost!
 - 2. **Just I[v2] write succeeds.**
 - No data block => will read garbage data from disk.
 - No bitmap entry, but inode has a pointer to Db => FS inconsistency!
 - 3. **Just B[v2] write succeeds.**
 - Bitmap says Db is allocated, inode has no pointer to it => again, FS inconsistent + Db can never be used again



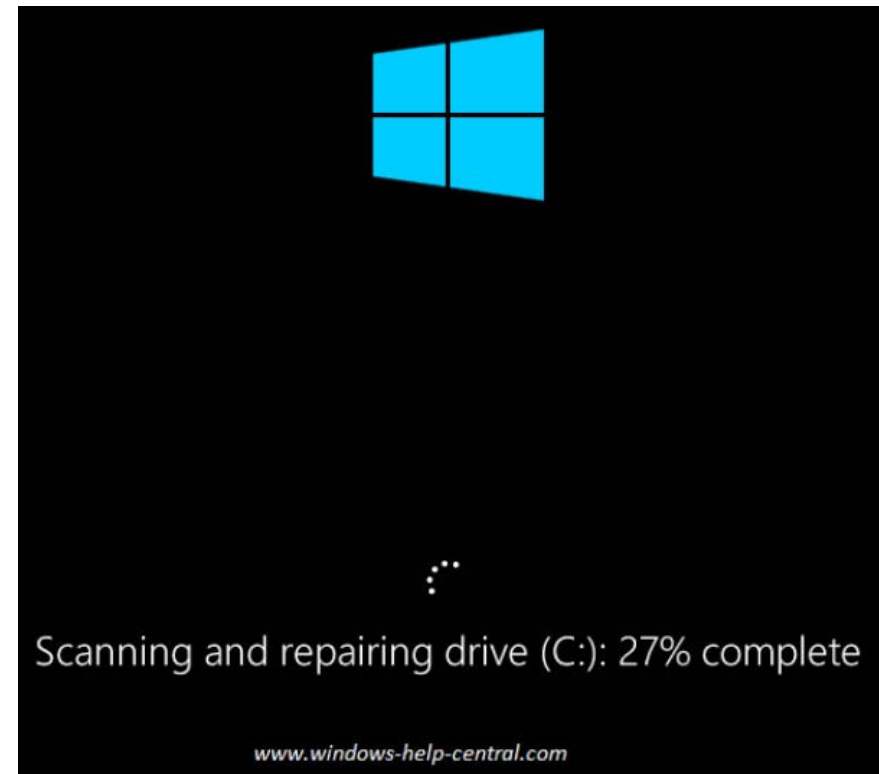
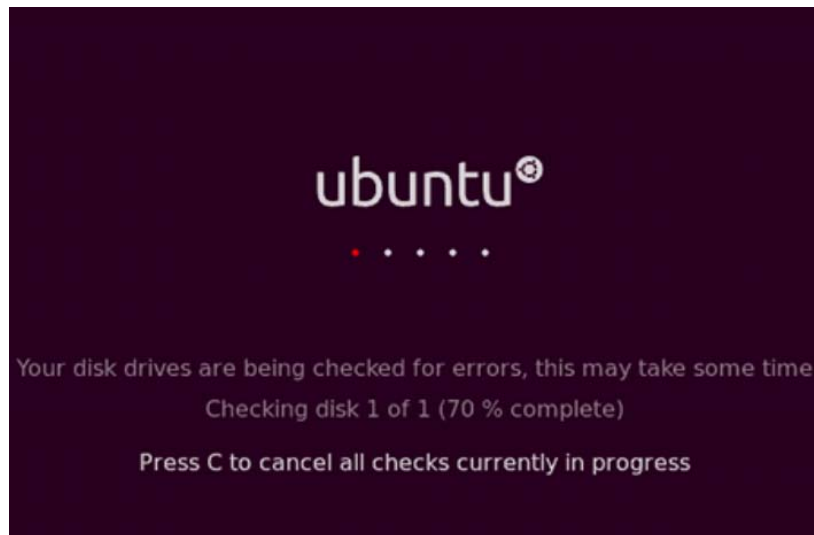
Other crash scenarios

- What if only two writes succeed before a crash?
 - 1. Only I[v2] and B[v2] writes succeed.
 - Inode and bitmap agree => FS metadata is consistent
 - However, Db contains garbage.
 - 2. Only I[v2] and Db writes succeed.
 - Inode points to correct data, but clashes with B[v1] => FS inconsistency, must fix this!
 - 3. Only B[v2] and Db writes succeed.
 - Again, inode and bitmap info does not match
 - Even though Db was written, no inode points to it (which file is it part of? No idea!)



Crash consistency problem

Solution #1: fsck



Similar tools exist on various systems



Solution: fsck

- fsck: UNIX tool for finding inconsistencies and repairing them
- Cannot fix all problems!
 - When Db is garbage – cannot know that's the case
 - Only cares that FS metadata is consistent!



Solution: fsck

- What does it check?
 - 1. **Superblock**: sanity checks
 - Use another superblock copy if suspected corruption
 - 2. **Free blocks**: scan inodes (incl. all indirect blocks), build bitmap
 - inodes / data bitmaps inconsistency => resolve by trusting inodes
 - Ensure inodes in use are marked in inode bitmaps
 - 3. **Inode state**: check inode fields for possible corruption
 - e.g., must have a valid “mode” field (file, dir, link, etc.)
 - If cannot fix => remove inode and update inode bitmap
 - 4. **Inode links**: verify links# for each inode
 - Traverse directory tree, compute expected links#, fix if needed
 - If inode discovered, but no dir refers to it => move to “lost+found”



Solution: fsck

- 5. **Duplicates**: check if two different inodes refer to same block
 - Clear one if obviously bad, or, give each inode its own copy of block
- 6. **Bad blocks**: bad pointers (outside of valid range)
 - Just remove the pointer from the inode or indirect block
- 7. **Directory checks**: integrity of directory structure
 - E.g., make sure that “.” and “..” are the first entries, each inode in a directory entry is allocated, no directory is linked more than once



fsck limitations

- So, fsck helps ensure integrity
- Only FS integrity, cannot do anything about lost data!
- Bigger problem: **too slow!**
 - Disks are very large nowadays – scanning all this could take hours!
 - Even for small inconsistency, must scan whole disk!



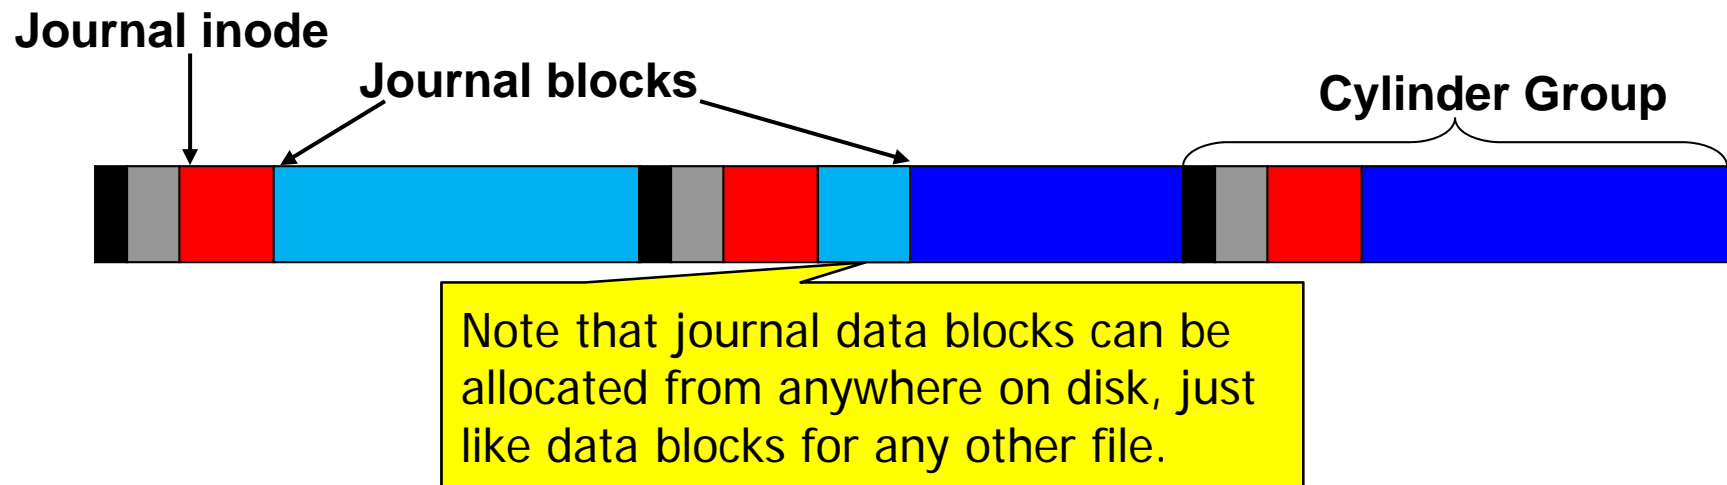
Alternative solution: Journaling

- Aka Write-Ahead-Logging
- Basic idea:
 - When doing an update, before overwriting structures, first write down a little note (elsewhere on disk) saying what you plan to do.
 - i.e., “Log” the operations you are about to do.
- If a crash takes place during the actual write => go back to journal and retry the actual writes.
 - Don't need to scan the entire disk, we know what to do!
 - Can recover data as well
- If a crash happens before journal write finishes, then it doesn't matter since the actual write has NOT happened at all, so nothing is inconsistent.



Linux Ext3 File System

- Extends ext2 with journaling capabilities
 - Backwards and forwards compatible
 - Identical on-disk format
 - Journal can be just another large file (inode, indirect blocks, data blocks)





What goes in that “note”

- Transaction structure:
 - Starts with a “transaction begin” (TxBegin) block, containing a transaction ID
 - Followed by blocks with the content to be written
 - Physical logging: log exact physical content
 - Logical logging: log more compact logical representation
 - Ends with a “transaction end” (TxEnd) block, containing the corresponding TID

Journal
entry

TxBegin (TID=1)	Updated inode	Updated Bitmap	Updated Data block	TxEnd (TID=1)
--------------------	------------------	-------------------	-----------------------	------------------



Data Journaling Example

- Say we have a regular update – add 1 data block to a file:
 - Write inode ($I[v2]$), Bitmap ($B[v2]$), Data block (Db)
 - Markers for the log (transaction begin/end)

TxBegin	$I[v2]$	$B[v2]$	Db	TxEnd
---------	---------	---------	----	-------



Data Journaling Example

- Say we have a regular update – add 1 data block to a file:
 - Write inode ($I[v2]$), Bitmap ($B[v2]$), Data block (Db)
 - Markers for the log (transaction begin/end)



- Sequence of operations
 - 1. Write the transaction (containing $Iv2$, $Bv2$, Db) to the log
 - 2. Write the blocks ($Iv2$, $Bv2$, Db) to the file system
 - 3. Mark the transaction free in the journal
- Crash may happen at any point!
 - If between 1 and 2 => on reboot, replay non-free transactions (called **redo logging**)
 - If during writes to the journal (step 1) => tricky!



Data Journaling Example

- One solution: write each block at a time
 - Slow!
 - Ideally issue multiple blocks at once.
 - Unsafe though! What could happen?
 - Normal operation: Blocks get written in order, power cuts off before TxEnd gets written => We know transaction is not valid, no problem.

TxBegin	I[v2]	B[v2]	Db	???
---------	-------	-------	----	-----

- However, Internal disk scheduling: TxBegin, Iv2, Bv2, TxEnd, Db
- Disk may lose power before Db written

TxBegin	I[v2]	B[v2]	???	TxEnd
---------	-------	-------	-----	-------

- Problem: Looks like a valid transaction!



Data journaling example

- To avoid this, split into 2 steps
 - 1. Write all except TxEnd to journal (**Journal Write step**)

TxBegin	I[v2]	B[v2]	Db
---------	-------	-------	----

- 2. Write TxEnd (only once 1. completes) (**Journal Commit step**)

=> final state is safe!

TxBegin	I[v2]	B[v2]	Db	TxEnd
---------	-------	-------	----	-------

- 3. Finally, now that journal entry is safe, write the actual data and metadata to their right locations on the FS (**Checkpoint step**)
- 4. Mark transaction as free in journal (**Free step**)



Journaling: Recovery Summary

- If crash happens before the transaction is committed to the journal
 - Just skip the pending update
- If crash happens during the checkpoint step
 - After reboot, scan the journal and look for committed transactions
 - Replay these transactions
 - After replay, the FS is guaranteed to be consistent
 - Called **redo logging**



Journal Space Requirements

- How much space do we need for the journal?
 - For every update, we log to the journal => sounds like it's huge!
- After “checkpoint” step, the transaction is not needed anymore because metadata and data made it safely to disk
 - So the space can be freed (free step).
- In practice: **circular log**.



Metadata Journaling

- Recovery is much faster with journaling
 - Replay only a few transactions instead of checking the whole disk
- However, normal operations are slower
 - Every update must write to the journal first, then do the update
 - Writing time is at least doubled
 - Journal writing may break sequential writing. Why?
 - Jump back-and-forth between writes to journal and writes to main region
 - Metadata journaling is similar, except we only write FS metadata (no actual data) to the journal:

Journal entry	TxBegin (TID=1)	Updated inode	Updated Bitmap	TxEnd (TID=1)
---------------	-----------------	---------------	----------------	---------------



Metadata Journaling

- What can happen now?
 - Say we write data after checkpointing metadata
 - If crash occurs before all data is written, inodes will point to garbage data!
 - How do we take care of this?
- Write data BEFORE writing metadata to journal!
 - 1. Write data, wait until it completes
 - 2. Metadata journal write
 - 3. Metadata journal commit
 - 4. Checkpoint metadata
 - 5. Free
- If write data fails => as if nothing happened, sort of (from the FS's point of view)!
- If write metadata fails => same!



Summary: Journaling

- Journaling ensures file system consistency
- Complexity is in the size of the journal, not the size of the disk!
- Is fsck useless then?
- Metadata journaling is the most commonly used
 - Reduces the amount of traffic to the journal, and provides reasonable consistency guarantees at the same time.
- Widely adopted in most modern file systems (ext3, ext4, ReiserFS, JFS, XFS, NTFS, etc.)



Ext3 final notes

- Lacks modern FS features (e.g., extents)
 - For recoverability, this may actually be an advantage
 - FS metadata is in fixed, well-known locations, and data structures have redundancy
 - When faced with significant data corruption, ext2/3 may be recoverable when a tree-based file-system may not



- Next up: Log-structured file systems