

CSC369 - Tutorial 8

Ext2 overview (metadata and files)

Starter code explained

Exercises 7 and 8 (no handout, they're online)


Ext2 overview

- Ext2 splits the HD into **blocks** of 1024 bytes. Think of it as an array of blocks!

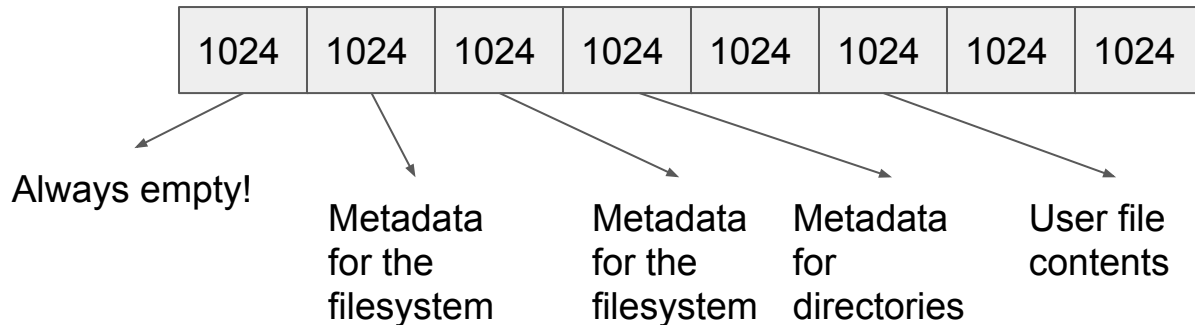
[illegible]

Ext2 overview

- Ext2 splits the HD into **blocks** of 1024 bytes. Think of it as an array of blocks!




- And assigns different semantics to each of them.



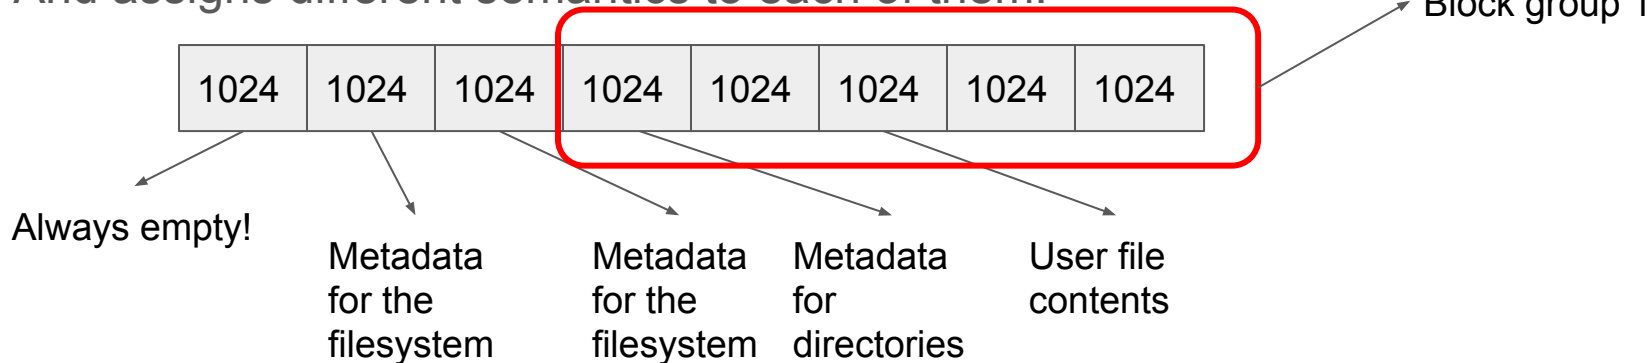
Don't use this order as reference, this is merely illustrative.

Ext2 overview

- Ext2 splits the HD into **blocks** of 1024 bytes. Think of it as an array of blocks!



- And assigns different semantics to each of them.



- Ext2 groups blocks into **block groups**
 - For A4, you can assume only one block group.

Ext2 overview

- Each file has an associated **inode**.
 - **This does not contain the file itself!**
- A data structure with information about the file! Which information?

```

/*
 * Structure of an inode on the disk
 */
struct ext2_inode {
    unsigned short i_mode;          /* File mode */
    /* Use 0 as the user id for the assignment. */
    unsigned short i_uid;          /* Low 16 bits of Owner Uid */
    unsigned int i_size;           /* Size in bytes */
    /* You don't need to set access time for the assignment. */
    unsigned int i_atime;          /* Access time */
    unsigned int i_ctime;          /* Creation time */
    /* You don't need to set modification time for the assignment. */
    unsigned int i_mtime;          /* Modification time */
    /* d_time must be set when appropriate */
    unsigned int i_dtime;          /* Deletion Time */
    /* Use 0 as the group id for the assignment. */
    unsigned short i_gid;          /* Low 16 bits of Group Id */
    unsigned short i_links_count; /* Links count */
    unsigned int i_blocks;         /* Blocks count IN DISK SECTORS*/
    /* You can ignore flags for the assignment. */
    unsigned int i_flags;          /* File flags */
    /* You should set it to 0. */
    unsigned int osd1;             /* OS dependent 1 */
    unsigned int i_block[15];      /* Pointers to blocks */
    /* You should use generation number 0 for the assignment. */
    unsigned int i_generation;     /* File version (for NFS) */
    /* The following fields should be 0 for the assignment. */
    unsigned int i_file_acl;       /* File ACL */
    unsigned int i_dir_acl;        /* Directory ACL */
    unsigned int i_faddr;          /* Fragment address */
    unsigned int extra[3];
};

```

```

/*
 * Structure of an inode on the disk
 */
struct ext2 inode {
    unsigned short i_mode;        /* File mode */
    /* Use 0 as the user id for the assignment. */
    unsigned short i_uid;         /* Low 16 bits of Owner */
    unsigned int i_size;          /* Size in bytes */
    /* You don't need to set access time for the assignment */
    unsigned int i_atime;         /* Access time */
    unsigned int i_ctime;         /* Creation time */
    /* You don't need to set modification time for the assignment */
    unsigned int i_mtime;        /* Modification time */
    /* d_time must be set when appropriate */
    unsigned int i_dtime;        /* Deletion Time */
    /* Use 0 as the group id for the assignment. */
    unsigned short i_gid;         /* Low 16 bits of Group Id */
    unsigned short i_links_count; /* Links count */
    unsigned int i_blocks;        /* Blocks count IN DISK SECTORS */
    /* You can ignore flags for the assignment. */
    unsigned int i_flags;         /* File flags */
    /* You should set it to 0. */
    unsigned int osd1;            /* OS dependent 1 */
    unsigned int i_block[15];     /* Pointers to blocks */
    /* You should use generation number 0 for the assignment. */
    unsigned int i_generation;    /* File version (for NFS) */
    /* The following fields should be 0 for the assignment. */
    unsigned int i_file_acl;      /* File ACL */
    unsigned int i_dir_acl;       /* Directory ACL */
    unsigned int i_faddr;         /* Fragment address */
    unsigned int extra[3];
};
*/
/*
 * Type field for file mode
 */
#define EXT2_S_IFLNK 0xA000 /* symbolic link */
#define EXT2_S_IFREG 0x8000 /* regular file */
#define EXT2_S_IFDIR 0x4000 /* directory */
/* Other types, irrelevant for the assignment */
/* #define EXT2_S_IFSOCK 0xC000 */ /* socket */
/* #define EXT2_S_IFBLK 0x6000 */ /* block device */
/* #define EXT2_S_IFCHR 0x2000 */ /* character device */
/* #define EXT2_S_IFIFO 0x1000 */ /* fifo */

```



```

/*
 * Structure of an inode on the disk
 */
struct ext2_inode {
    unsigned short i_mode;          /* File mode */
    /* Use 0 as the user id for the assignment. */
    unsigned short i_uid;          /* Low 16 bits of Owner Uid */
    unsigned int i_size;           /* Size in bytes */
    /* You don't need to set access time for the assignment. */
    unsigned int i_atime;          /* Access time */
    unsigned int i_ctime;          /* Creation time */
    /* You don't need to set modification time for the assignment. */
    unsigned int i_mtime;          /* Modification time */
    /* d_time must be set when appropriate */
    unsigned int i_dtime;          /* Deletion Time */
    /* Use 0 as the group id for the assignment. */
    unsigned short i_gid;          /* Low 16 bits of Group Id */
    unsigned short i_links_count; /* Links count */
    unsigned int i_blocks;         /* Blocks count IN DISK SECTORS*/
    /* You can ignore flags for the assignment. */
    unsigned int i_flags;          /* File flags */
    /* You should set it to 0. */
    unsigned int osd1;             /* OS dependent 1 */
    unsigned int i_block[15];      /* Pointers to blocks */
    /* You should use generation number 0 for the assignment. */
    unsigned int i_generation;     /* File version (for NFS) */
    /* The following fields should be 0 for the assignment. */
    unsigned int i_file_acl;       /* File ACL */
    unsigned int i_dir_acl;        /* Directory ACL */
    unsigned int i_faddr;          /* Fragment address */
    unsigned int extra[3];
};

```



```

/*
 * Structure of an inode on the disk
 */
struct ext2_inode {
    unsigned short i_mode;          /* File mode */
    /* Use 0 as the user id for the assignment. */
    unsigned short i_uid;          /* Low 16 bits of Owner Uid */
    unsigned int i_size;           /* Size in bytes */
    /* You don't need to set access time for the assignment. */
    unsigned int i_atime;          /* Access time */
    unsigned int i_ctime;          /* Creation time */
    /* You don't need to set modification time for the assignment. */
    unsigned int i_mtime;          /* Modification time */
    /* d_time must be set when appropriate */
    unsigned int i_dtime;          /* Deletion Time */
    /* Use 0 as the group id for the assignment. */
    unsigned short i_gid;          /* Low 16 bits of Group Id */
    unsigned short i_links_count; /* Links count */
    unsigned int i_blocks;         /* Blocks count IN DISK SECTORS */
    /* You can ignore flags for the assignment. */
    unsigned int i_flags;          /* File flags */
    /* You should set it to 0. */
    unsigned int osd1;             /* OS dependent 1 */
    unsigned int i_block[15];      /* Pointers to blocks */
    /* You should use generation number 0 for the assignment. */
    unsigned int i_generation;     /* File version (for NFS) */
    /* The following fields should be 0 for the assignment. */
    unsigned int i_file_acl;       /* File ACL */
    unsigned int i_dir_acl;       /* Directory ACL */
    unsigned int i_faddr;          /* Fragment address */
    unsigned int extra[3];
};

```

What is this? There are important details you **must** know in order to complete this assignment.

```

/*
 * Structure of an inode on the disk
 */
struct ext2_inode {
    unsigned short i_mode;          /* File mode */
    /* Use 0 as the user id for the assignment. */
    unsigned short i_uid;          /* Low 16 bits of Owner Uid */
    unsigned int i_size;           /* Size in bytes */
    /* You don't need to set access time for the assignment. */
    unsigned int i_atime;          /* Access time */
    unsigned int i_ctime;          /* Creation time */
    /* You don't need to set modification time for the assignment. */
    unsigned int i_mtime;          /* Modification time */
    /* d_time must be set when appropriate */
    unsigned int i_dtime;          /* Deletion Time */
    /* Use 0 as the group id for the assignment. */
    unsigned short i_gid;          /* Low 16 bits of Group Id */
    unsigned short i_links_count; /* Links count */
    unsigned int i_blocks;         /* Blocks count IN DISK SECTORS*/
    /* You can ignore flags for the assignment. */
    unsigned int i_flags;          /* File flags */
    /* You should set it to 0. */
    unsigned int osd1;             /* OS dependent 1 */
    unsigned int i_block[15];      /* Pointers to blocks */
    /* You should use generation number 0 for the assignment. */
    unsigned int i_generation;     /* File version (for NFS) */
    /* The following fields should be 0 for the assignment. */
    unsigned int i_file_acl;       /* File ACL */
    unsigned int i_dir_acl;        /* Directory ACL */
    unsigned int i_faddr;          /* Fragment address */
    unsigned int extra[3];
};

```

What is this? There are important details you **must** know in order to complete this assignment.

Check the documentation!!

<http://www.nongnu.org/ext2-doc/ext2.html#I-LINKS-COUNT>

```

/*
 * Structure of an inode on the disk
 */
struct ext2_inode {
    unsigned short i_mode;          /* File mode */
    /* Use 0 as the user id for the assignment. */
    unsigned short i_uid;          /* Low 16 bits of Owner Uid */
    unsigned int i_size;           /* Size in bytes */
    /* You don't need to set access time for the assignment. */
    unsigned int i_atime;          /* Access time */
    unsigned int i_ctime;          /* Creation time */
    /* You don't need to set modification time for the assignment. */
    unsigned int i_mtime;          /* Modification time */
    /* d_time must be set when appropriate */
    unsigned int i_dtime;          /* Deletion Time */
    /* Use 0 as the group id for the assignment. */
    unsigned short i_gid;          /* Low 16 bits of Group Id */
    unsigned short i_links_count; /* Links count */
    unsigned int i_blocks;         /* Blocks count IN DISK SECTORS*/
    /* You can ignore flags for the assignment. */
    unsigned int i_flags;          /* File flags */
    /* You should set it to 0. */
    unsigned int osd1;             /* OS dependent 1 */
    unsigned int i_block[15];      /* Pointers to blocks */
    /* You should use generation number 0 for the assignment. */
    unsigned int i_generation;     /* File version (for NFS) */
    /* The following fields should be 0 for the assignment. */
    unsigned int i_file_acl;       /* File ACL */
    unsigned int i_dir_acl;       /* Directory ACL */
    unsigned int i_faddr;          /* Fragment address */
    unsigned int extra[3];
};

```

Another pitfall here... check the documentation!


```

/*
 * Structure of an inode on the disk
 */
struct ext2_inode {
    unsigned short i_mode;          /* File mode */
    /* Use 0 as the user id for the assignment. */
    unsigned short i_uid;          /* Low 16 bits of Owner Uid */
    unsigned int i_size;           /* Size in bytes */
    /* You don't need to set access time for the assignment. */
    unsigned int i_atime;          /* Access time */
    unsigned int i_ctime;          /* Creation time */
    /* You don't need to set modification time for the assignment. */
    unsigned int i_mtime;          /* Modification time */
    /* d_time must be set when appropriate */
    unsigned int i_dtime;          /* Deletion Time */
    /* Use 0 as the group id for the assignment. */
    unsigned short i_gid;          /* Low 16 bits of Group Id */
    unsigned short i_links_count; /* Links count */
    unsigned int i_blocks;         /* Blocks count IN DISK SECTORS*/
    /* You can ignore flags for the assignment. */
    unsigned int i_flags;          /* File flags */
    /* You should set it to 0. */
    unsigned int osd1;             /* OS dependent 1 */
    unsigned int i_block[15];      /* Pointers to blocks */
    /* You should use generation number 0 for the assignment. */
    unsigned int i_generation;     /* File version (for NFS) */
    /* The following fields should be 0 for the assignment. */
    unsigned int i_file_acl;       /* File ACL */
    unsigned int i_dir_acl;        /* Directory ACL */
    unsigned int i_faddr;          /* Fragment address */
    unsigned int extra[3];
};

```

The most important part of the inode!
An array of block numbers, telling you in which blocks the **contents** of the file are located!

```

/*
 * Structure of an inode on the disk
 */
struct ext2_inode {
    unsigned short i_mode;          /* File mode */
    /* Use 0 as the user id for the assignment. */
    unsigned short i_uid;          /* Low 16 bits of Owner Uid */
    unsigned int i_size;           /* Size in bytes */
    /* You don't need to set access time for the assignment. */
    unsigned int i_atime;          /* Access time */
    unsigned int i_ctime;          /* Creation time */
    /* You don't need to set modification time for the assignment. */
    unsigned int i_mtime;          /* Modification time */
    /* d_time must be set when appropriate */
    unsigned int i_dtime;          /* Deletion Time */
    /* Use 0 as the group id for the assignment. */
    unsigned short i_gid;          /* Low 16 bits of Group Id */
    unsigned short i_links_count; /* Links count */
    unsigned int i_blocks;         /* Blocks count IN DISK SECTORS*/
    /* You can ignore flags for the assignment. */
    unsigned int i_flags;          /* File flags */
    /* You should set it to 0. */
    unsigned int osd1;             /* OS dependent 1 */
    unsigned int i_block[15];      /* Pointers to blocks */
    /* You should use generation number 0 for the assignment. */
    unsigned int i_generation;     /* File version (for NFS) */
    /* The following fields should be 0 for the assignment. */
    unsigned int i_file_acl;       /* File ACL */
    unsigned int i_dir_acl;       /* Directory ACL */
    unsigned int i_faddr;         /* Fragment address */
    unsigned int extra[3];
};

```

If this is a text file with contents “Hello world”,
and $i_block[0] = 10$,
then the text is located in the 10th block,
byte number $10 \times 1024 = 10240$.

```

/*
 * Structure of an inode on the disk
 */
struct ext2_inode {
    unsigned short i_mode;          /* File mode */
    /* Use 0 as the user id for the assignment. */
    unsigned short i_uid;          /* Low 16 bits of Owner Uid */
    unsigned int i_size;           /* Size in bytes */
    /* You don't need to set access time for the assignment. */
    unsigned int i_atime;          /* Access time */
    unsigned int i_ctime;          /* Creation time */
    /* You don't need to set modification time for the assignment. */
    unsigned int i_mtime;          /* Modification time */
    /* d_time must be set when appropriate */
    unsigned int i_dtime;          /* Deletion Time */
    /* Use 0 as the group id for the assignment. */
    unsigned short i_gid;          /* Low 16 bits of Group Id */
    unsigned short i_links_count; /* Links count */
    unsigned int i_blocks;         /* Blocks count IN DISK SECTORS*/
    /* You can ignore flags for the assignment. */
    unsigned int i_flags;          /* File flags */
    /* You should set it to 0. */
    unsigned int osd1;             /* OS dependent 1 */
    unsigned int i_block[15];      /* Pointers to blocks */
    /* You should use generation number 0 for the assignment. */
    unsigned int i_generation;     /* File version (for NFS) */
    /* The following fields should be 0 for the assignment. */
    unsigned int i_file_acl;       /* File ACL */
    unsigned int i_dir_acl;        /* Directory ACL */
    unsigned int i_faddr;          /* Fragment address */
    unsigned int extra[3];
};

```

Read the documentation, there is much more to this than what I've just described.

Ext2 overview

- Each file has an associated **inode**.
 - **This is does not contain the file itself!**
- A data structure with information about the file! Which information?
- For each block group, Ext2 will allocate a constant number of consecutive blocks to serve as an array of inodes: the **inode table**.
 - (notice each inode has a fixed size)

Ext2 overview

- Each file has an associated **inode**.
 - **This is does not contain the file itself!**
- A data structure with information about the file! Which information?
- For each block group, Ext2 will allocate a constant number of consecutive blocks to serve as an array of inodes: the **inode table**.
 - (notice each inode has a fixed size)
 - In A4, what does that mean for the number of inodes available?

Ext2 overview

- Each file has an associated **inode**.
 - **This does not contain the file itself!**
- A data structure with information about the file! Which information?
- For each block group, Ext2 will allocate a constant number of consecutive blocks to serve as an array of inodes: the **inode table**.
 - (notice each inode has a fixed size)
 - In A4, what does that mean for the number of inodes available?
 - You have a very limited number of inodes, and therefore a maximum number of files you can store!
 - Some inodes are reserved... documentation!

Ext2 overview

- Suppose you have a pointer to your disk:

```
char *disk; //Suppose, for now, that your disk is in memory.
```

Ext2 overview

- Suppose you have a pointer to your disk:

`char *disk; //Suppose, for now, that your disk is in memory.`

- Suppose your inode table is at block 10,

`struct ext2_inode *inodes = (struct ext2_inode *)(disk + 1024 * 10);`

Ext2 overview

- Suppose you have a pointer to your disk:

```
char *disk; //Suppose, for now, that your disk is in memory.
```

- Suppose your inode table is at block 10,

```
struct ext2_inode *inodes = (struct ext2_inode *)(disk + 1024 * 10);
```

- and you want to print the contents of the 7th inode...

```
printf("Size in bytes = %d", inodes[6].i_size);
```

Ext2 overview

- Suppose Ext2 has 32 inodes in total. When you create a file, you need grab a free inode to store the file's metadata.

Ext2 overview

- Suppose Ext2 has 32 inodes in total. When you create a file, you need grab a free inode to store the file's metadata.
- How do we find a free inode?
 - Iterate over the inode table? But the inode structure doesn't have a member indicating whether it is associated with a file... It could be filled with garbage values...

Ext2 overview

- Suppose Ext2 has 32 inodes in total. When you create a file, you need grab a free inode to store the file's metadata.
- How do we find a free inode?
 - Iterate over the inode table? But the inode structure doesn't have a member indicating whether it is associated with a file... It could be filled with garbage values...
- Ext2 allocates a block to be the **inode bitmap**!
 - A vector of bits, where `inode_bitmap[i]` = 1 iff the i-th inode is in use.

Ext2 overview

- Suppose your inode bitmap is located at block 4.

Ext2 overview

- Suppose your inode bitmap is located at block 4.

```
unsigned char *inode_bits = (unsigned char *)(disk + 1024 * 4);
```

Ext2 overview

- Suppose your inode bitmap is located at block 4.

```
unsigned char *inode_bits = (unsigned char *)(disk + 1024 * 4);
```

- From the documentation:
 - *The first inode of this block group is represented by bit 0 of byte 0, the second by bit 1 of byte 0. The 8th inode is represented by bit 7 (most significant bit) of byte 0 while the 9th inode is represented by bit 0 (least significant bit) of byte 1.*

Ext2 overview

- Suppose your inode bitmap is located at block 4.

```
unsigned char *inode_bits = (unsigned char *)(disk + 1024 * 4);
```

- From the documentation:
 - *The first inode of this block group is represented by bit 0 of byte 0, the second by bit 1 of byte 0. The 8th inode is represented by bit 7 (most significant bit) of byte 0 while the 9th inode is represented by bit 0 (least significant bit) of byte 1.*
- If you have 32 inodes, you are looking at the first 32 bits, i.e., the first 4 bytes of the bitmap:

```
for (byte = 0; byte < 32 / 8; byte++)  
    for (bit = 0; bit < 8; bit++)  
        in_use = inode_bits[byte] & (1 << bit);
```

Ext2 overview

- Recall that file contents are also stored in blocks.
- So when you add a new file, you also need find free blocks for it.

Ext2 overview

- Recall that file contents are also stored in blocks.
- So when you add a new file, you also need find free blocks for it.
- Just like for inodes, Ext2 will allocate a special block to be the **block bitmap!**

Ext2 overview

- How do we find:
 - The inode table?
 - The block bitmap?
 - The inode bitmap?

Ext2 overview

- How do we find:
 - The inode table?
 - The block bitmap?
 - The inode bitmap?

Ext2 overview

- How do we find:
 - The inode table?
 - The block bitmap?
 - The inode bitmap?
- Ext2 has a special block called **block group descriptor table**.
 - An array of **block group descriptors**

Ext2 overview

- How do we find:
 - The inode table?
 - The block bitmap?
 - The inode bitmap?
- Ext2 has a special block called **block group descriptor table**.
 - An array of **block group descriptors**

Ext2 overview

- How do we find:
 - The inode table?
 - The block bitmap?
 - The inode bitmap?
- Ext2 has a special block called **block group descriptor table**.
 - An array of **block group descriptors**
 - Recall that there is only one block group in the disks used in A4.

Ext2 overview

```
/*
 * Structure of a blocks group descriptor
 */
struct ext2_group_desc
{
    unsigned int    bg_block_bitmap;    /* Blocks bitmap block */
    unsigned int    bg_inode_bitmap;    /* Inodes bitmap block */
    unsigned int    bg_inode_table;     /* Inodes table block */
    unsigned short  bg_free_blocks_count; /* Free blocks count */
    unsigned short  bg_free_inodes_count; /* Free inodes count */
    unsigned short  bg_used_dirs_count; /* Directories count */
    /* The pad and reserved fields should be 0 for the assignment. */
    unsigned short  bg_pad;
    unsigned int    bg_reserved[3];
};
```

Ext2 overview

- How do we find:
 - The inode table?
 - The block bitmap?
 - The inode bitmap?
- Ext2 has a special block called **block group descriptor table**.
 - An array of **block group descriptors**
 - Recall that there is only one block group in the disks used in A4.
- How do we find this block?
 - It is, for this assignment, the third block in the disk.

Ext2 overview

- There is one last important data structure: the **superblock**.
- It contains information about the whole filesystem:
 - number available inodes/blocks across all groups,
 - constants related to sizes, etc.
 - the code is too big to paste here, but it is the **ext2_super_block** struct in the header.

Ext2 overview

- What does your typical A4 disk image actually look like:

Ext2 overview

- What does your typical A4 disk image actually look like:




1024

Empty space


Ext2 overview

- What does your typical A4 disk image actually look like:



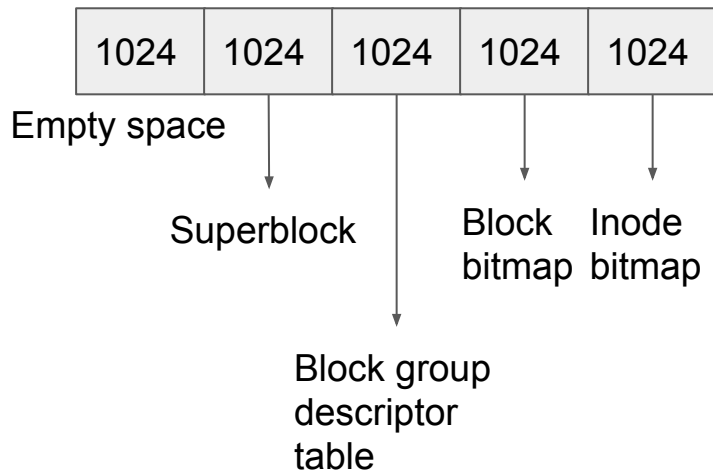
Ext2 overview

- What does your typical A4 disk image actually look like:



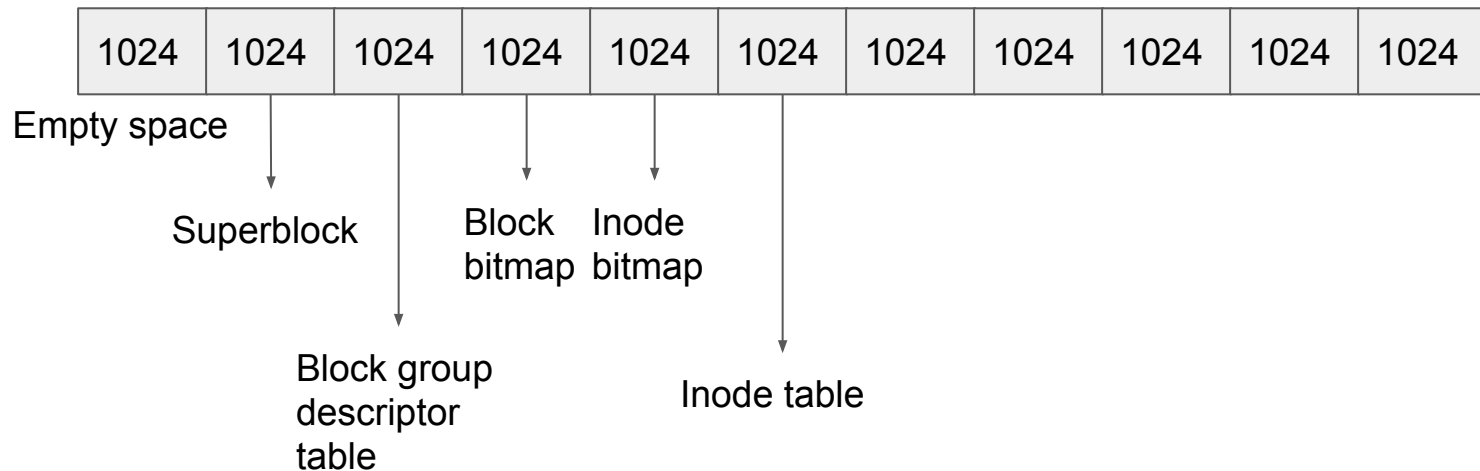
Ext2 overview

- What does your typical A4 disk image actually look like:



Ext2 overview

- What does your typical A4 disk image actually look like:



char *disk;

Starter code

```
13 int main(int argc, char **argv) {
14
15     if(argc != 2) {
16         fprintf(stderr, "Usage: %s <image file name>\n", argv[0]);
17         exit(1);
18     }
19     int fd = open(argv[1], O_RDWR);
20
21     disk = mmap(NULL, 128 * 1024, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
22     if(disk == MAP_FAILED) {
23         perror("mmap");
24         exit(1);
25     }
26
27     struct ext2_super_block *sb = (struct ext2_super_block *) (disk + 1024);
28     printf("Inodes: %d\n", sb->s_inodes_count);
29     printf("Blocks: %d\n", sb->s_blocks_count);
30
31     return 0;
32 }
```

Exercises 7 and 8

- Print some attributes of the block group descriptor.
 - Exercise 7
- Print some attributes of each inode that is being used.
 - Exercise 8
 - Some inodes are reserved and should be ignored. Which ones? Look at the documentation.

Looking at dumps: debugging your code

- It will be useful to examine the images you create/modify using external tools.
- Dump the files emptydisk.img and onefile.img
 - `xxd -c 1 emptydisk.img emptydisk.dump`
 - `xxd -c 1 onefile.img onefile.dump`
- This dumps one byte per line. You can change this, but I find it useful to have line numbers matching the byte number you have.

Looking at dumps: debugging your code

- It will be useful to examine the images you create/modify using external tools.
- Dump the files `emptydisk.img` and `onefile.img`
 - `xxd -c 1 emptydisk.img emptydisk.dump`
 - `xxd -c 1 onefile.img onefile.dump`
- This dumps one byte per line. You can change this, but I find it useful to have line numbers matching the byte number you have.
- Now open your files with:
 - `vimdiff emptydisk.dump onefile.dump`

Looking at dumps: debugging your code

free blocks count

free inodes count

Mind the endianness!

```
1031 00000406: 00 .
1032 00000407: 00 .
1033 00000408: 06 .
1034 00000409: 00 .
1035 0000040a: 00 .
1036 0000040b: 00 .
1037 0000040c: 69 i
1038 0000040d: 00 .
1039 0000040e: 00 .
1040 0000040f: 00 .
1041 00000410: 15 .
1042 00000411: 00 .
1043 00000412: 00 .
1044 00000413: 00 .
1045 00000414: 01 .
1046 00000415: 00 .
1047 00000416: 00 .
```

```
1031 00000406: 00 .
1032 00000407: 00 .
1033 00000408: 06 .
1034 00000409: 00 .
1035 0000040a: 00 .
1036 0000040b: 00 .
1037 0000040c: 68 h
1038 0000040d: 00 .
1039 0000040e: 00 .
1040 0000040f: 00 .
1041 00000410: 14 .
1042 00000411: 00 .
1043 00000412: 00 .
1044 00000413: 00 .
1045 00000414: 01 .
1046 00000415: 00 .
1047 00000416: 00 .
```

Vim uses 1 based indexing for lines, so we're one byte off (you can simply delete the first line to fix that...)

Directories

- A directory is represented just like a file:
 - An inode + some data blocks.
- While the data blocks for a file store its contents, a directory data block stores a linked list of **directory entries**.
- **An entry must ALWAYS have a length that is a multiple of 4.**
 - **The last entry must always occupy the remaining space of the block.**

Directories

```
struct ext2_dir_entry {  
    unsigned int    inode;        /* Inode number */  
    unsigned short  rec_len;      /* Directory entry length */  
    unsigned char   name_len;     /* Name length */  
    unsigned char   file_type;  
    char            name[];       /* File name, up to EXT2_NAME_LEN */  
};
```

- **An entry must ALWAYS have a length that is a multiple of 4.**
 - The last entry must always occupy the remaining space of the block.
- This is why we have two “length” fields:
 - To jump to the following record, just do `base + rec_len` (how do you tell when you have reached the last entry?).
 - To print the name, we use `name_len`.

Directories

```
struct ext2_dir_entry {  
    unsigned int    inode;        /* Inode number */  
    unsigned short  rec_len;      /* Directory entry length */  
    unsigned char   name_len;     /* Name length */  
    unsigned char   file_type;  
    char            name[];       /* File name, up to EXT2_NAME_LEN */  
};
```

- Suppose some directory stores its contents in block 9.

Directories

```
struct ext2_dir_entry {  
    unsigned int    inode;        /* Inode number */  
    unsigned short  rec_len;      /* Directory entry length */  
    unsigned char    name_len;    /* Name length */  
    unsigned char    file_type;  
    char             name[];      /* File name, up to EXT2_NAME_LEN */  
};
```

- Suppose some directory stores its contents in block 9.
- How do you find the name of the first entry?
 - `char* name = disk + 1024 * 9 + sizeof(ext2_dir_entry)`
 - ... problems?

Directories

```
struct ext2_dir_entry {  
    unsigned int    inode;        /* Inode number */  
    unsigned short  rec_len;      /* Directory entry length */  
    unsigned char   name_len;     /* Name length */  
    unsigned char   file_type;  
    char            name[];       /* File name, up to EXT2_NAME_LEN */  
};
```

- Suppose some directory stores its contents in block 9.
- How do you find the name of the first entry?
 - `char* name = disk + 1024 * 9 + sizeof(ext2_dir_entry)`
 - ... problems?
 - Isn't this "eating" the first 4 bytes (or however many bytes a pointer needs) of the name?

Directories

```
struct ext2_dir_entry {  
    unsigned int    inode;        /* Inode number */  
    unsigned short  rec_len;      /* Directory entry length */  
    unsigned char   name_len;     /* Name length */  
    unsigned char   file_type;  
    char            name[];       /* File name, up to EXT2_NAME_LEN */  
};
```

- Suppose some directory stores its contents in block 9.
- How do you find the name of the first entry?
 - `char* name = disk + 1024 * 9 + sizeof(ext2_dir_entry)`
 - Isn't this "eating" the first 4 bytes (or however many bytes a pointer needs) of the name?
 - No! In C, declaring `char name[]`; is an indication that the struct will have variable size. That member has "size" 0!

Directories

```
struct ext2_dir_entry {  
    unsigned int    inode;        /* Inode number */  
    unsigned short  rec_len;      /* Directory entry length */  
    unsigned char   name_len;     /* Name length */  
    unsigned char   file_type;  
    char            name[];       /* File name, up to EXT2_NAME_LEN */  
};
```

- `char* name = disk + 1024 * 9 + sizeof(ext2_dir_entry)`
- How do we print this? This is **not** NULL terminated.

Directories

```
struct ext2_dir_entry {  
    unsigned int    inode;        /* Inode number */  
    unsigned short  rec_len;      /* Directory entry length */  
    unsigned char    name_len;    /* Name length */  
    unsigned char    file_type;  
    char             name[];      /* File name, up to EXT2_NAME_LEN */  
};
```

- `char* name = disk + 1024 * 9 + sizeof(ext2_dir_entry)`
- How do we print this? This is **not** NULL terminated.
 - `ext2_dir_entry *entry = (ext2_dir_entry*) disk + 1024* 9;`
 - `printf("%.s", entry->name_len, name);`

Directories

```
struct ext2_dir_entry {  
    unsigned int    inode;        /* Inode number */  
    unsigned short  rec_len;      /* Directory entry length */  
    unsigned char   name_len;     /* Name length */  
    unsigned char   file_type;  
    char            name[];       /* File name, up to EXT2_NAME_LEN */  
};
```

- `char* name = disk + 1024 * 9 + sizeof(ext2_dir_entry)`
- How do we print this? This is **not** NULL terminated.
 - `ext2_dir_entry *entry = (ext2_dir_entry*) disk + 1024* 9;`
 - `printf("%.s", entry->name_len, name);`
- The next entry is located at...?

Directories

```
struct ext2_dir_entry {  
    unsigned int    inode;        /* Inode number */  
    unsigned short  rec_len;      /* Directory entry length */  
    unsigned char    name_len;    /* Name length */  
    unsigned char    file_type;  
    char             name[];      /* File name, up to EXT2_NAME_LEN */  
};
```

- `char* name = disk + 1024 * 9 + sizeof(ext2_dir_entry)`
- How do we print this? This is **not** NULL terminated.
 - `ext2_dir_entry *entry = (ext2_dir_entry*) disk + 1024* 9;`
 - `printf("%.s", entry->name_len, name);`
- The next entry is located at...?
 - `((char*) entry)+ entry->rec_len;`

Exercise 9

- Iterate over all inodes.
- If the inode is being used, and if it is related to a directory,
 - Access the directory data blocks and print each directory entry.
- Some pitfalls:
 - How do you detect deleted entries? Do you need to?
 - What if there is a deleted entry as the very first entry of the block?
 - Again... check the documentation for answers (or ask me next week).