

CSC369 Tutorial 4

Synchronization Primitives

(continued)

Synchronization Mechanisms

- Locks
 - Very primitive constructs with minimal semantics
- Semaphores
 - A generalization of locks
 - Easy to understand, hard to program with
- Condition Variables
 - Constructs used in implementing *monitors* (more on this later)

Semaphores

- Synchronization mechanism that generalizes locks to more than just “acquired” and “free” (or “released”)
- A semaphore provides you with:
 - An integer counter accessed through 2 atomic operations
 - Wait - aka: down, decrement, P (for proberen)
 - Block until semaphore is free (counter > 1), then decrement the counter
 - Signal - aka: up, post, increment, V (for verhogen)
 - Increment the counter and unblock one waiting thread (if there are any)
 - A queue of waiting threads
- A mutex is just a binary semaphore
 - remember `pthread_mutex_lock()` blocks if another thread is holding the lock

POSIX Semaphores

- Declared in semaphore.h
- A few calls associated with POSIX semaphores:
 - sem_init
 - Initialize the semaphore
 - sem_wait
 - Wait on the semaphore (decrement value)
 - sem_post
 - Signal (post) on the semaphore (increment value)
 - sem_getvalue
 - Get the current value of the semaphore
 - sem_destroy
 - Destroy the semaphore

Initializing & Destroying POSIX Semaphores

- Initialize semaphores using `sem_init`
`int sem_init(sem_t *sem, int pshared, unsigned int value);`
 - `sem`: the semaphore to initialize
 - `pshared`: non-zero to share between processes (e.g. after fork)
 - `value`: initial count value of the semaphore
- Destroy semaphores using `sem_destroy`
`int sem_destroy(sem_t *sem);`
 - `sem`: semaphore to destroy
 - Semaphore must have been created using `sem_init`
 - Destroying a semaphore that has threads blocked on it is undefined
- No static initializer. Why?

Decrementing & Incrementing POSIX Semaphores

- Decrement semaphores using `sem_wait`
`int sem_wait(sem_t *sem);`
 - `sem`: the semaphore to decrement (wait on)
- Increment semaphores using `sem_post`
`int sem_post(sem_t *sem);`
 - `sem`: semaphore to increment
- Let's look at an example of a very simple server simulation

Semaphores: (main) use case

- Allow only up to N threads to access a resource at any point in time

```
sem_init(&semaphore, N);
```

```
...
```

```
sem_wait(&semaphore);
```

```
// Only up to N threads can be executing here
```

```
sem_post(&semaphore);
```

- Example: server accepting $\leq N$ clients simultaneously
- After N threads have executed `sem_wait()`, thread # $N+1$ will block until one of them releases the semaphore

Server Example

```
//...
#define NUM_THREADS 200
#define NUM_RESOURCES 10
sem_t resource_sem; // Semaphore declaration

int main (int argc, char *argv[]) {
    pthread_t thread[NUM_THREADS];
    sem_init(&resource_sem, 0, NUM_RESOURCES); // Resource Semaphore

    for (int i = 0; i < NUM_THREADS; i++) {
        int rc = pthread_create(&thread[i], NULL, handle_connection, (void*)i);
        if (rc != 0) {
            printf("ERROR: pthread_create() returned %d\n", rc);
            exit(-1);
        }
    }
    //...
    for (int i = 0; i < NUM_THREADS; i++) {
        void *status = NULL;
        int rc = pthread_join(thread[i], &status);
        if (rc != 0) {
            printf("ERROR: pthread_join() returned %d\n", rc);
            exit(-1);
        }
    }
    return 0;
} // End of main
```


Server Example – Connection Handler

```
void *handle_connection(void *c)
{
    int client = (int)c;
    printf("Handler for client %d created!\n", client);

    sem_wait(&resource_sem);

    // DO WORK TO HANDLE CONNECTION HERE
    sleep(1);
    printf("Done servicing client %d\n", client);

    sem_post(&resource_sem);

    return NULL;
}
```

Condition Variables

- Another useful synchronization construct used in implementing monitors; only a single process executes inside the monitor
- Locks control thread access to data; condition variables allow threads to synchronize based on the value of the data
 - a thread can notify another thread that a condition is now satisfied
- Alternative to condition variables is to constantly poll the data (from the critical section); it's a bad idea:
 - Ties up a lot of CPU resources
 - Could potentially lead to synchronization problems
- Conditional variable operations:
 - `wait()` – suspend the invoking process and release the lock
 - `signal()` – resume exactly one suspended process
 - `broadcast()` – resumes all suspended processes
 - If no process is suspended, `signal/broadcast` has no effect (in contrast to semaphores, where `signal` always changes state of the semaphore)

POSIX Condition Variables

- POSIX condition variables: `pthread_cond_t`

- A few calls associated with POSIX CVs:

```
int pthread_cond_init(pthread_cond_t *cond, pthread_condattr_t  
*attr);
```

- Initialize a condition variable
- Also **PTHREAD_COND_INITIALIZER**

```
int pthread_cond_destroy(pthread_cond_t *cond);
```

- Destroy a condition variable

```
int pthread_cond_wait(pthread_cond_t *cond,  
pthread_mutex_t *mutex);
```

- Wait on a condition variable; *mutex* is released during wait

```
int pthread_cond_signal(pthread_cond_t *cond);
```

- Wake up one thread waiting on this condition variable

```
int pthread_cond_broadcast(pthread_cond_t *cond);
```

- Wake up all threads waiting on this condition variable

POSIX Condition Variables

- `int pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex);`
 - Two steps:
 1. **Atomically** release *mutex* and start waiting for *cond* to be signaled
 2. When *cond* is signaled by another thread, **atomically** acquire *mutex* and stop waiting, continue execution
- Signaling thread must release the mutex after calling *pthread_cond_signal()* / *pthread_cond_broadcast()* before *pthread_cond_wait()* can return

Using Condition Variables (from LLNL tutorial)

Main Thread

- Declare and initialize global data/variables which require synchronization (such as "count")
- Declare and initialize a condition variable object
- Declare and initialize an associated mutex
- Create threads A and B to do work

Thread A

- Do work up to the point where a certain condition must occur (such as "count" must reach a specified value)
- Lock associated mutex and check value of a global variable
- Call `pthread_cond_wait()` to perform a blocking wait for signal from Thread-B. Note that a call to `pthread_cond_wait()` automatically and atomically unlocks the associated mutex variable so that it can be used by Thread-B.
- When signalled, wake up. Mutex is automatically and atomically locked.
- Explicitly unlock mutex
- Continue

Thread B

- Do work
- Lock associated mutex
- Change the value of the global variable that Thread-A is waiting upon.
- Check value of the global Thread-A wait variable. If it fulfills the desired condition, signal Thread-A.
- Unlock mutex.
- Continue

Main Thread: Join / Continue

Monitors

- Locks
 - Provide mutual exclusion
 - 2 operations: `acquire()` and `release()`
- Semaphores
 - Generalize locks with an integer count variable and a thread queue
 - 2 operations: `wait()` and `signal()`
 - If the integer count is negative, threads wait in a queue until another thread signals the semaphore
- Monitors
 - An abstraction that encapsulates shared data and operations on it in such a way that only a single process at a time may be executing “in” the monitor

More on Monitors

- Programmer defines the scope of the monitor
 - ie: which data is “monitored”
 - basically, “monitor” == “synchronized object”
- Local data can be accessed only by the monitor’s procedures (not by any external procedures)
- Before any monitor procedure may be invoked, mutual exclusion must be guaranteed
 - There is often a lock associated with each monitored object
- Other processes that attempt to enter the monitor are blocked. They must first acquire the lock before becoming active in the monitor

Complications With Monitors

- Complication
 - A process may need to wait for something to happen
 - Input from another thread might be necessary for example
 - The other thread may require access to the monitor to produce that event
- Solution?
 - Monitors support suspending execution within the monitor
 - wait() – suspend the invoking process and release the lock
 - signal() – resume exactly one suspended process
 - broadcast() – resumes all suspended processes
 - If no process is suspended, signal/broadcast has no effect (in contrast to semaphores, where signal always changes state of the semaphore)

Monitor signal(); who goes first?

- Suppose P executes a signal operation that would wake up a suspended process Q
 - Either process can continue execution, but both cannot simultaneously be active in the monitor
- Who goes first?
 - Hoare monitors: waiter first
 - signal() immediately switches from the caller to a waiting thread
 - Condition that the waiter was blocked on is guaranteed to hold when the waiter resumes
 - Mesa monitors: signaler first
 - signal() places a waiter on the ready queue, but signaler continues inside the monitor
 - Condition that the waiter was blocked on is not guaranteed to hold when the waiter resumes (must check again...)

Hoare vs. Mesa Monitors

- Hoare monitor wait

```
if (...) {  
    wait(cv, lock);  
}
```
- Mesa monitor wait

```
while (...) {  
    wait(cv, lock);  
}
```
- Tradeoffs
 - Hoare monitors are easier to reason with, but hard to implement
 - Mesa monitors are easier to implement, and support additional operations like broadcast()
 - Pthread conditional variables implement Mesa semantics

Monitors implementation

- Implemented (e.g.) using pthreads conditional variables and mutexes
- One mutex (to protect access to data)
- Possibly multiple condition variables
 - one per each condition to be signaled/waited on
- Note: mutex necessary for pthread conditional variables even if not used for monitors
 - mutex protects the conditional variable itself;
race between waiting thread and signaling thread
 - signal/broadcast do nothing if no waiting threads

Monitor Example - Bounded Buffers

- We have a buffer of limited size N
 - Producers add to the buffer if it is not full
 - Consumers remove from the buffer if it is not empty
- Want to control buffer as a monitor
 - Buffer can only be accessed by methods that are “part of” the monitor, that only give one producer or consumer access to the buffer at a time
- Need 2 functions
 - `add_to_buffer()`
 - `remove_from_buffer()`
- Need
 - One lock
 - Two conditions
 - One for producers to wait
 - One for consumers to wait

Monitor Example - Bounded Buffers

```
#define N 100
```

```
typedef struct buf_s {  
    int data[N];  
    int inpos; // producer inserts here  
    int outpos; // consumer removes from here  
    int numelements; // # of items in buffer  
    struct lock *lock; // access to monitor  
    struct cv *notFull; // for producers to wait  
    struct cv *notEmpty; // for consumers to wait  
} buf_t;
```

```
buf_t buffer;  
void add_to_buff(int value);  
int remove_from_buff();
```


Monitor Example - Bounded Buffers

```
void add_to_buf(int value)
{
    lock_acquire(buffer.lock);

    while (buffer.numelements == N) {
        // buffer is full, wait
        cv_wait(buffer.notFull, buffer.lock);
    }

    buffer.data[buffer.inpos] = value;
    buffer.inpos = (buffer.inpos + 1) % N;
    buffer.numelements++;

    cv_signal(buffer.notEmpty);
    lock_release(buffer.lock);
}
```



What kind of
monitor is this?

Monitor Example - Bounded Buffers

```
int remove_from_buf()
{
    lock_acquire(buffer.lock);

    while (buffer.numelements == 0) {
        // buffer is empty, wait
        cv_wait(buffer.notEmpty, buffer.lock);
    }

    int val = buffer.data[buffer.outpos];
    buffer.outpos = (buffer.outpos + 1) % N;
    buffer.numelements--;

    cv_signal(buffer.notFull);
    lock_release(buffer.lock);
    return val;
}
```