

# CSC369 Tutorial 2

Kernel modules  
Synchronization (basics)  
A1 & system calls


**Handout is being distributed,  
read it and you'll find the answers  
along today's tutorial.**

**Vote on the poll for A1 office hours!**

# Kernel

- Kernel = computer program that connects the user applications to the system hardware
- Handles:
  - Memory management
  - CPU scheduling (Process and task management)
  - Disk management
  - User access to other I/O devices (e.g., network card)

# Kernel



Source: wikipedia.org

# Kernel modules

- Object file that contains code to extend the kernel's functionality
- Why do we need them? Why not include all possible functionality in the kernel directly?
  - Kernel code lies in main memory (limited resource)
  - Kernel should be minimal
  - Avoid functionality bloating
  - For each new functionality added => recompile kernel, reboot, ... ugh!
  - Instead, develop modules separately, load as needed
  - Modularity => Better chance to recover from buggy new code without a complete kernel crash!

# Why should I bother?

- Because it's cool!



- Better understanding on how the OS works
- Write awesome extensions to the OS
- Write your own device drivers!

# Linux kernel modules

- Basic utilities:
  - `insmod`: to load a module
  - `rmmmod`: to unload a module
- The `modprobe` utility
  - More complex, deals with module dependencies
  - We won't be using it
- Module objects: `.ko` files

# Loading/unloading a kernel module

- As root, or using sudo
- Loading:
  - `insmod mymodule.ko`
- Unloading:
  - `rmmod mymodule.ko` (or: `rmmod mymodule`)
- Entry point:
  - `module_init(mymodule_init);`
  - `module_exit(mymodule_exit);`

# Kernel module example

```
#include <linux/kernel.h>
#include <linux/init.h>
#include <linux/module.h>

MODULE_DESCRIPTION("My kernel module");
MODULE_AUTHOR("John Doe");
MODULE_LICENSE("GPL");

static int mymodule_init(void) {
    printk( KERN_DEBUG "Hello world!\n" );
    return 0;
}

static void mymodule_exit(void) {
    printk( KERN_DEBUG "I'm outta here\n" );
}

module_init(mymodule_init);
module_exit(mymodule_exit);
```



# Compiling a module

- Different than a regular C program
- Must use different headers
- Must not link with libraries. Why?
- Must be compiled with the same options as the kernel in which we want to load it
- Standard method: kbuild
  - Two files: a Makefile, and a Kbuild file

# Example

- **Makefile:**

```
KDIR=/lib/modules/`uname -r`/build
```

```
kbuild:
```

```
    make -C $(KDIR) M=`pwd`
```

```
clean:
```

```
    make -C $(KDIR) M=`pwd` clean
```

- **Kbuild file:**

```
EXTRA_CFLAGS=-g
```

```
obj-m = mymodule.o
```

# Printing messages

- Use printk
  - e.g., `printk( KERN_DEBUG "Hello world\n" );`
- Dude, where's my output?
  - Not displayed at stdout
  - Can be retrieved from the system logs
  - Use `dmesg` command
    - can use `grep` to search and filter:  
`dmesg | grep Hello -C 3`

# Debugging a kernel module

- More complicated than a regular program
- A bug in a module can lead to the whole OS malfunctioning
- Buggy module: can lead to a “kernel oops”
- Avoid reboot cycles => use VM for CSC369!
  - VM snapshots can be useful
- Do not develop modules directly on your Linux box without a VM! – painfully slow!
- For A1, use rudimentary (yet efficient) method: `printk` statements

# Debugging a kernel module

- You can use a debugger, but not very useful
  - Simple bugs can be tracked easily with printks
  - Use ksymoops utility
- Complex bugs – not even a debugger will help as much
  - Need to know in depth the OS structure
  - Multiple contexts, interrupts, VM, etc.
- Kernel oops message can be translated using ksymoops (memory locations, backtrace, etc.)

–<http://opensourceforu.com/2011/01/understanding-a-kernel-oops/>

# Linux kernel API – some differences

- Different headers – make sure to include them!
- Success/failure conventions:
  - 0 == success
  - Non-zero (*negative*) == failure (-ENOMEM, -EINVAL, etc.)
  - See `<include/asm-generic/errno-base.h>` and `<include/asm-generic/errno.h>`

- Memory allocation: `kmalloc/kfree`

```
#include <linux/malloc.h>
if(!(string = kmalloc(len+1, GFP_KERNEL))) {
    return -ENOMEM;
}
```

...

```
kfree(string);
```

- GFP\_KERNEL can lead to suspending the current process; so it cannot be used in an interrupt context.

- GFP\_ATOMIC when the `kmalloc` won't suspend the current process; can be used anywhere

# Strings and printing

- Standard string functions:
  - strcmp, str(n)cpy, str(n)cat, memcpy, etc.
- Same header: <string.h>
- Printing: printk, defined in <linux/kernel.h>
- Similar syntax, plus category of message:  

```
printk(KERN_WARNING "Uh-oh, you better check this: %s\n", buff);  
printk(KERN_DEBUG "This buffer looks spooky: %s\n", buff);  
#define KERN_EMERG "<0>" /* system is unusable */  
#define KERN_ALERT "<1>" /* action must be taken immediately */  
#define KERN_CRIT "<2>" /* critical conditions */  
#define KERN_ERR "<3>" /* error conditions */  
#define KERN_WARNING "<4>" /* warning conditions */  
#define KERN_NOTICE "<5>" /* normal but significant condition */  
#define KERN_INFO "<6>" /* informational */  
#define KERN_DEBUG "<7>" /* debug-level messages */
```

# Synchronization

- Concurrent multiple threads accessing the same resource (e.g. shared data)
- Things can go wrong if access is not protected
  - example: 2 threads incrementing a counter:
    - increment is *not atomic*, reading the counter and writing the value back are separate operations

*// Thread 1:*

```
updated = counter + 1;
```

```
counter = updated;
```

*// Thread 2:*

```
updated = counter + 1;
```

```
counter = updated;
```



# Synchronization: locks

- Synchronization mechanisms with 2 operations: acquire (lock), and release (unlock)
- In simplest terms: an object associated with a particular critical section that you need to “own” if you wish to execute in that region
- Simple semantics to provide mutual exclusion:  
    acquire(lock);  
        // CRITICAL SECTION  
    release(lock);
- Downsides:
  - Can cause deadlock if not careful
  - Cannot allow multiple concurrent accesses to a resource

# Synchronization: spinlocks

- Busy-waiting synchronization
- Based on hardware atomic instructions, e.g. TAS – test-and-set
  - pseudocode:

```
bool test_and_set(bool *lock) {  
    // Atomically set to true and  
    // return the old value  
    bool old = *lock;  
    *lock = true;  
    return old;  
}
```

Suppose many threads are calling this function with the same pointer, Which threads will get False as a return value?

# Synchronization: spinlocks

- TAS spinlock implementation:

```
bool lock = false;
```

```
void acquire(bool *lock) {  
    // Wait until the lock becomes available  
    while (test_and_set(lock));  
}
```

```
void release(bool *lock) {  
    // Make the lock available  
    *lock = false;  
}
```

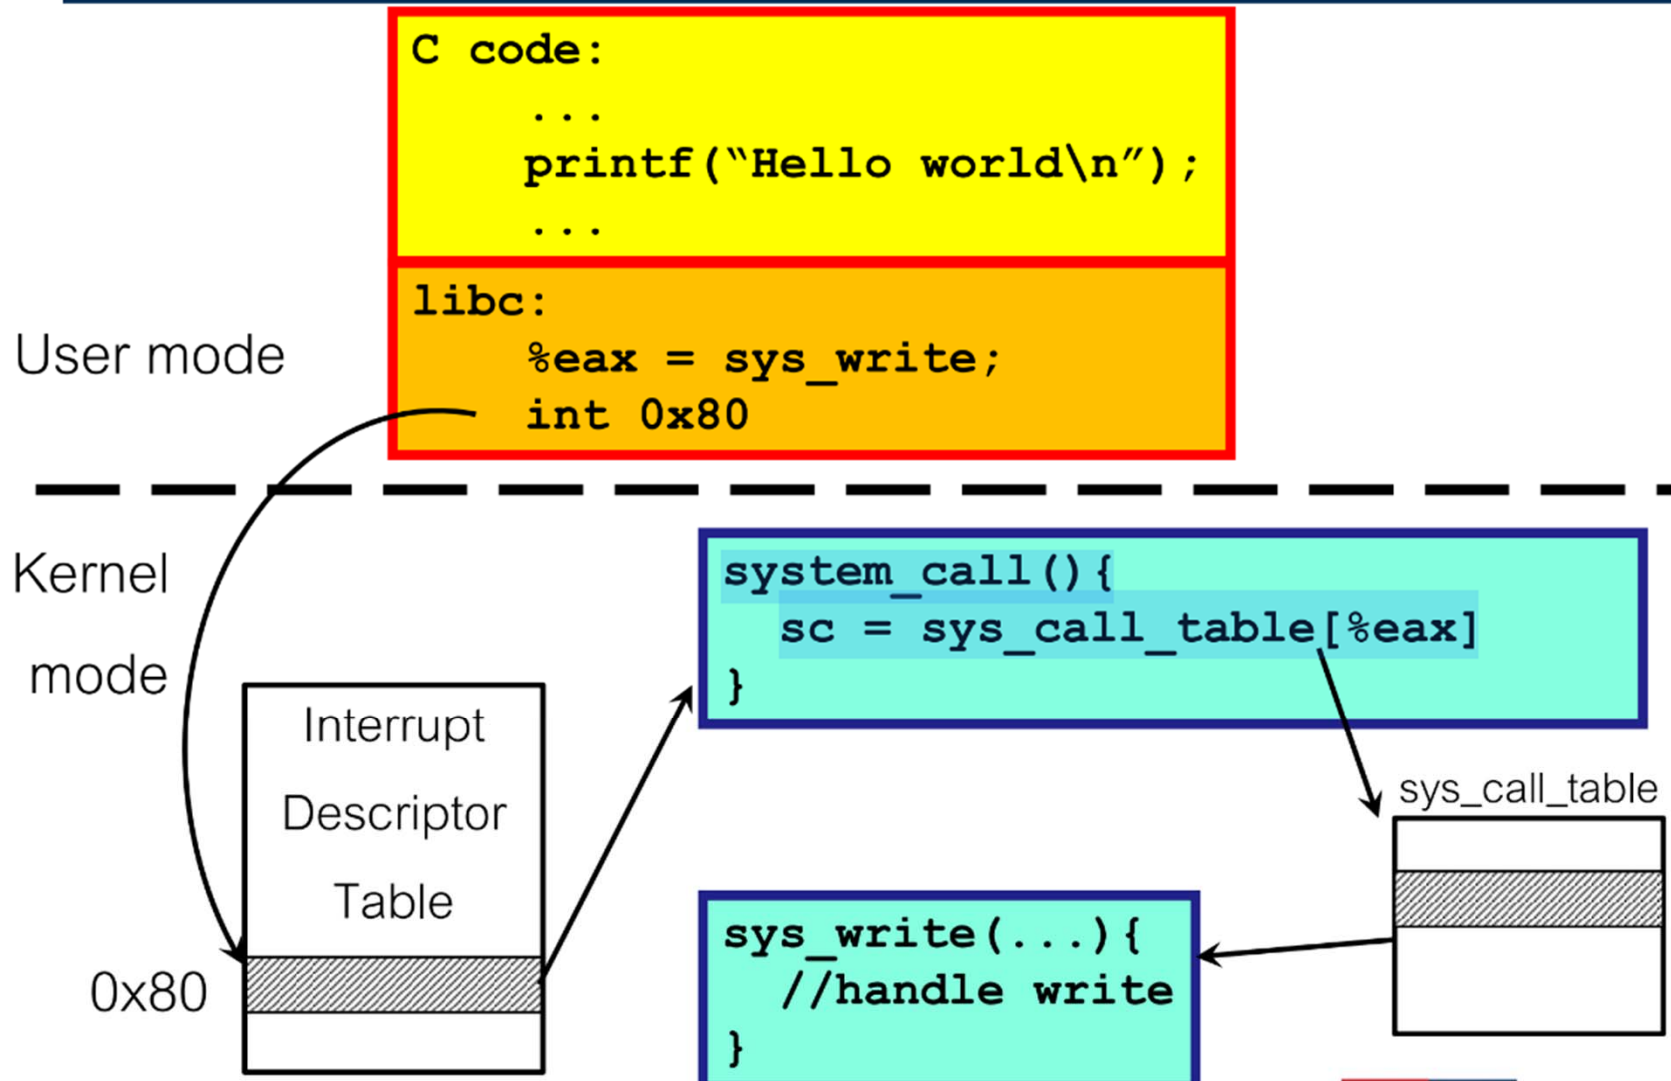
# Synchronization: spinlocks

- Linux kernel spinlocks: `spinlock_t` type
  - `spinlock_t myspinlock = SPIN_LOCK_UNLOCKED;`
- Operations:
  - `spin_lock_init(&myspinlock)`
  - `spin_lock/unlock(&myspinlock)`
- Can also use read/write spinlocks: `rwlock_t`
  - `rwlock_init()`, `read_lock()`, `write_lock()`
- Check out: `<include/linux/spinlock.h>`
- Will learn more about synchronization later in the course!

# System calls

- Application:  
`write(fd, buffer, size);`
- libc wrapper for `write()`:  
`eax = __NR_write;`  
`ebx = buffer;`  
`ecx = size;`  
*//...(up to 6 parameters in total)*  
`sysenter;`  
`return;`
- Kernel syscall trap handler:  
`syscall_table[eax](ebx, ecx, ...);`  
`sysexit;`

# A normal system call



# System calls: fun fact

- Which syscall is the most frequent?

# System calls: fun fact

- Which syscall is the most frequent?

Answer: *gettimeofday()*

(at least *one of* the most frequent)

Linux even has special optimizations for this call  
(a memory page mapped into both the kernel  
space and the user space for every process)



# Assignment 1

- Monitoring *system calls* made by processes
  - “highjack” syscalls of interest: substitute entries in the syscall table with our own function that records call info and invokes the original
  - Some system calls: open, read, write, close, wait, exec, fork, exit, and kill
- Kernel programming
- Test and debug in a virtual machine
- Read the handout and comments in the starter code very thoroughly

# VM setup for A1

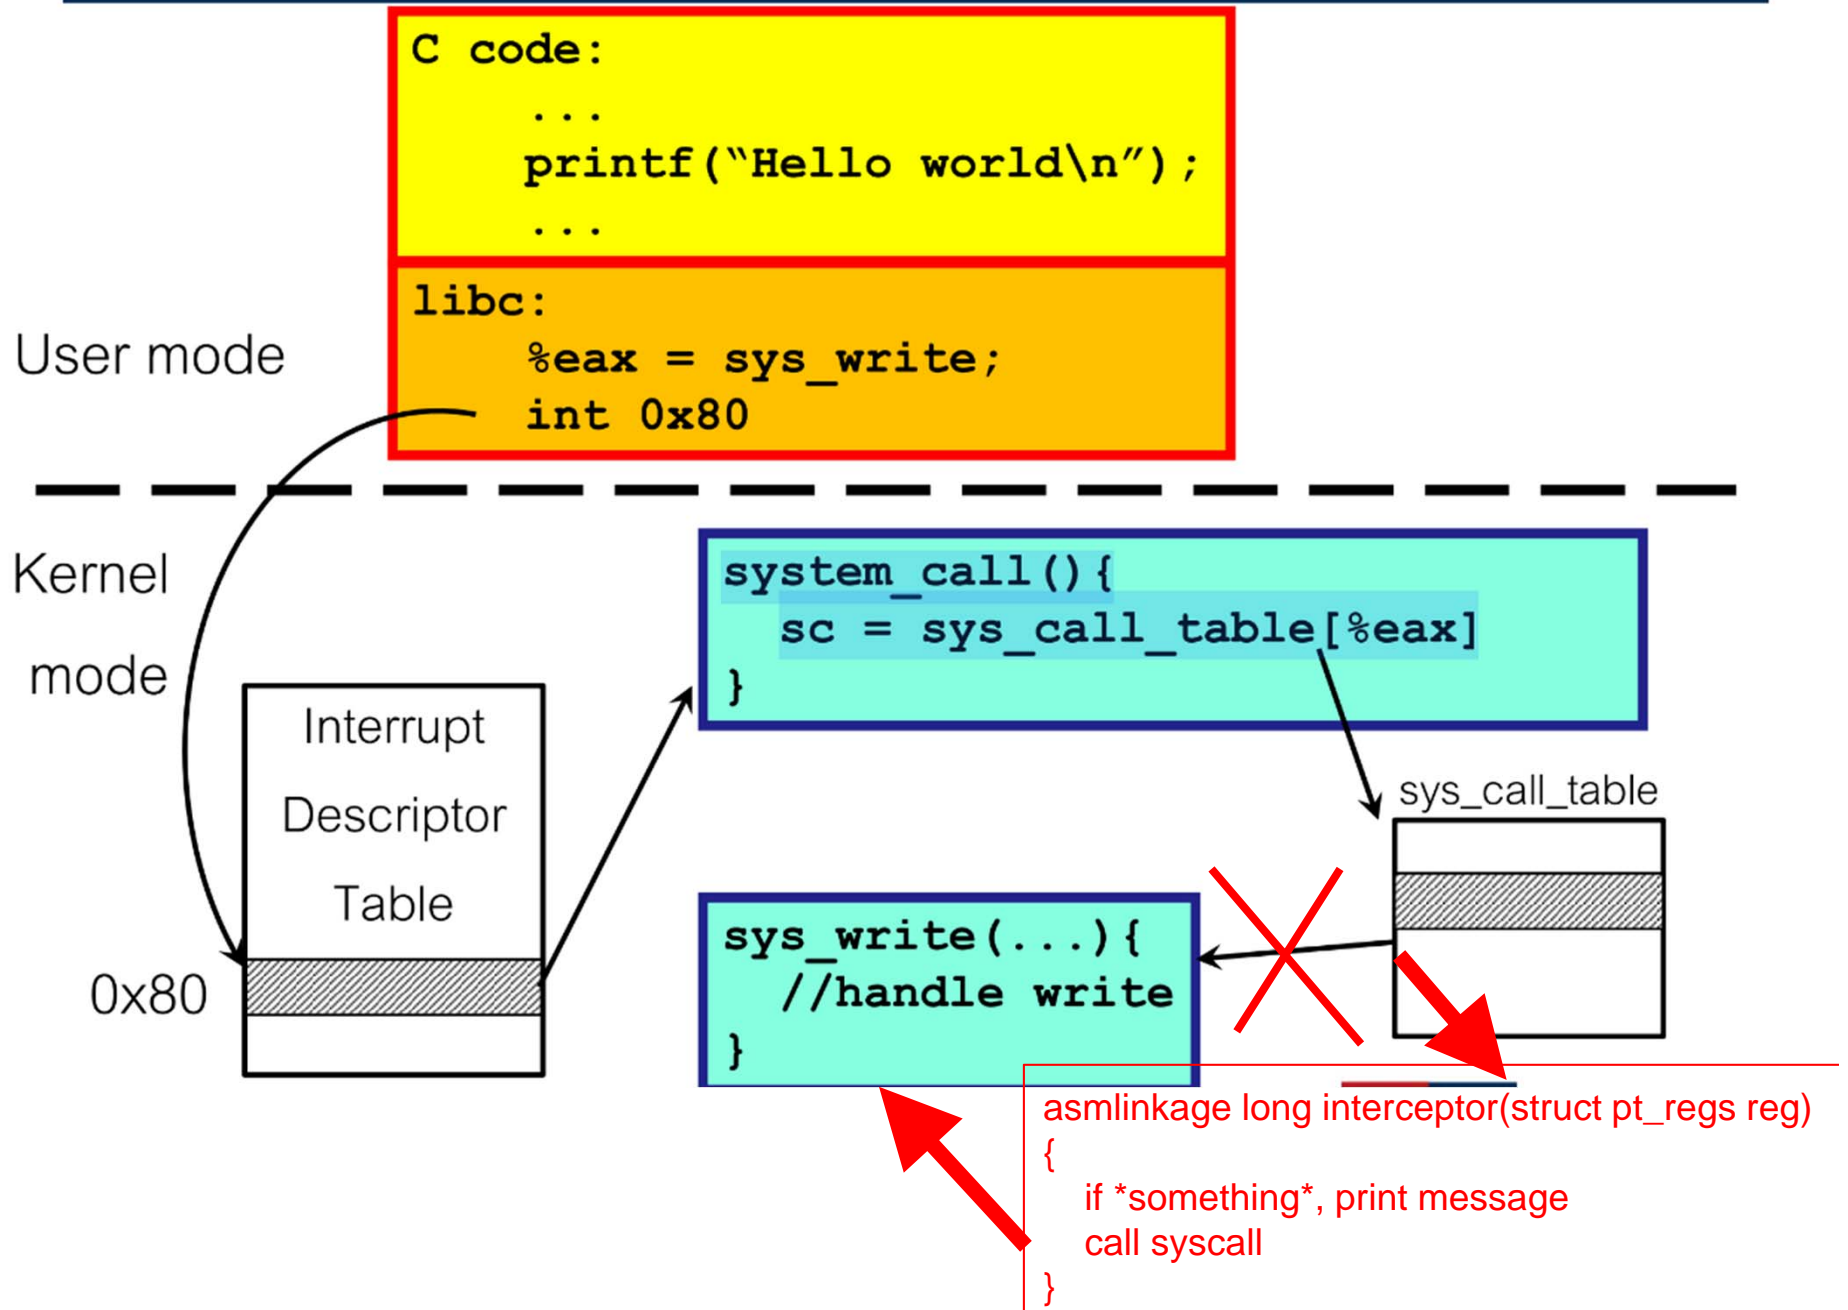
- Don't test your kernel module on a physical machine, use a virtual machine (VirtualBox/VMWare/etc.)
- A .vmdk virtual disk file is provided
- Use CDF lab machines or your own computers; follow instructions in the assignment handout
- Develop on physical machine, test in VM

# VM setup for A1

- Ways to transfer files to the VM:
  - Shared folders; need Guest Additions to be installed in the VM
  - Enable network in the VM, setup ssh server, then scp/rsync/svn checkout/etc.
    - also ssh into the VM (to have a nicer terminal)

Check Piazza, assignment handout, and google for instructions and help with VM setup

# An intercepted system call



# Some A1 tips

- Be careful with corner cases;  
document any assumptions that you make
  - as comments in code, or in your info.txt file
- Cleanup everything properly on unload; need to restore ***all*** kernel state that you have modified
- Passing all the provided tests doesn't guarantee that everything works correctly
- Be very careful with synchronization
  - e.g. do not forget to release locks before return when handling error conditions